A Ledger-Aware Benchmark for Market Infrastructure under Frequent Batch Auctions

Author Name
Institution
email@example.com

March 7, 2026

Abstract

We present a benchmark-oriented extension of a ledger-first market infrastructure prototype for studying market-design and execution behavior under settlement and risk constraints. The environment combines seedable agent-based order flow, configurable immediate, speed-bump, and frequent-batch-auction matching regimes, double-entry style account state transitions, and deterministic safety checks for conservation and non-negativity. In addition to throughput, latency, spread, price impact, queue-priority advantage, and arbitrage-profit proxies, the benchmark now reports welfare and behavior signals including retail surplus per traded unit, retail adverse-selection rate, welfare dispersion, and surplus-transfer gap. The current version exposes a step-wise `Reset/Step/Observe/Metrics` API, an adapter-driven control surface, four policy baselines including an offline contextual controller, and a unified four-dimensional stress sweep over arbitrage intensity, retail intensity, informed intensity, and maker quote width. Across all published runs, settlement invariants remain intact while mechanism and controller choices induce clear latency, fairness, and welfare tradeoffs.

1 Introduction

Electronic markets are shaped jointly by matching rules, latency structure, and settlement semantics. Market-design work on frequent batch auctions argues that short batching windows can reduce some latency-driven distortions in continuous markets [2]. In practice, many evaluation artifacts study matching rules in isolation and do not model the settlement layer, replay behavior, or safety checks that an actual market stack must satisfy.

This paper develops a benchmark-oriented layer on top of an existing ledger-first market-infrastructure prototype. It is not a replacement for the original systems paper. It is a separate evaluation track that asks a narrower question: can market-design and controller tradeoffs be measured inside an environment that preserves explicit settlement invariants and exposes richer welfare signals?

The current paper makes five concrete contributions:

1. It defines a ledger-aware benchmark environment in which market-design experiments are coupled to deterministic settlement checks.

2. It provides a seedable multi-agent order-flow generator with four agent classes, thirteen benchmark scenarios, explicit agent/workload sweeps, and a unified four-dimensional stress surface over arbitrage, retail, informed, and maker-width multipliers.

3. It reports single-seed and eight-seed aggregate benchmark results over throughput, latency, spread, price impact, queue-priority advantage, arbitrage-profit proxy, and welfare/behavior metrics.

4. It exposes a step-wise `Reset/Step/Observe/Metrics` API, a gym-style adapter with five runtime control channels, and four adapter-driven controllers, including a stronger offline contextual baseline.

5. It preserves executable settlement invariants across all artifacts, so controller improvements are evaluated under non-negativity and conservation constraints rather than in a matching-only simulator.

2 Problem Setting

We study a discrete-time environment in which a market state is advanced one simulator step at a time. At step t, the state is represented as

$$s_t = (B_t, S_t, X_t, f_t),$$

where B_t and S_t denote the buy and sell books, X_t is the vector of account states, and f_t is the latent synthetic fundamental used by the order-flow generator. A scenario additionally fixes a matching mode, a batch-window size, risk thresholds, and a population of agents.

The benchmark focuses on four concrete questions:

1. How do immediate, speed-bump, adaptive, and frequent-batch-auction execution differ in latency and fill behavior?

2. Can fairness-adjacent signals such as queue-priority advantage and latency-arbitrage profit be measured reproducibly under deterministic settlement semantics?

3. Do welfare/behavior metrics such as retail surplus and adverse selection reveal controller trade-offs that proxy metrics alone miss?

4. Do settlement invariants hold under heterogeneous agents and stress scenarios?

3 Benchmark Environment

3.1 Core Components

The benchmark environment lives in the repository's `simulator/` package. It consists of scenario, order, fill, and result schemas, heterogeneous agent order generation, immediate, speed-bump, adaptive, and batch-clearing matching paths, deterministic settlement application, metric accumulation, and benchmark artifact writers.

The environment also exposes a step-wise control surface with `Reset()`, `Observe()`, `Step()`, and `Metrics()`. On top of that surface, the repository includes an adapter layer that returns reward-bearing timesteps and advertises per-scenario action support. The current implementation exposes five runtime control channels: batch-window override for adaptive scenarios, risk-limit scaling, tie-break randomization, release-cadence control, and price-aggression bias [1].

3.2 Determinism and Invariants

Each scenario is seeded. Given the same scenario configuration and seed, the environment is expected to produce the same accepted order counts, fill counts, and benchmark metrics. The benchmark currently enforces three operational invariants:

1. **Conservation**: total cash and total inventory across accounts remain constant over time.
2. **Non-negativity**: no account is allowed to end a fill application with negative cash or negative units.
3. **Deterministic replay**: seeded runs produce stable outputs for the same configuration.

4 Agent Models

The present environment includes four agent classes:

- **Market makers**: quote on both sides around the current fundamental with configurable width and intensity.
- **Latency arbitrageurs**: react to stale price information and attempt to exploit short-horizon dislocations.
- **Retail traders**: submit noisier buy or sell orders with wider dispersion around the current reference level.
- **Informed traders**: trade in the direction of the next-step fundamental when the generator exposes a directional signal.

These agents are intentionally simple. Their purpose is to generate heterogeneous but controlled activity so that matching rules and controller choices can be compared under a repeatable workload.

5 Tasks and Baselines

The current benchmark suite contains thirteen scenarios:

1. `Immediate-Surrogate`
2. `SpeedBump-50ms`
3. `FBA-100ms`
4. `FBA-250ms`
5. `FBA-500ms`
6. `Adaptive-100-250ms`
7. `Adaptive-OrderFlow-100-250ms`
8. `Adaptive-QueueLoad-100-250ms`
9. `Policy-BurstAware-100-250ms`

10. `Policy-LearnedLinUCB-100-250ms`
11. `Policy-LearnedTinyMLP-100-250ms`
12. `Policy-LearnedOfflineContextual-100-250ms`
13. `FBA-250ms-Stress`

The controller suite now contains a hand-written burst-aware controller, a contextual linear bandit, a gradient-trained TinyMLP policy, and an offline contextual value controller trained from logged trajectories generated by burst-aware, LinUCB, TinyMLP, and random behavior policies.

6 Metrics

The environment reports five groups of signals:

- **Systems metrics**: orders per second, fills per second, and p50/p95/p99 latency.
- **Market-quality metrics**: average spread and average price impact.
- **Fairness-adjacent proxies**: queue-priority advantage, latency-arbitrage profit, and execution dispersion.
- **Welfare/behavior metrics**: retail surplus per traded unit, arbitrageur surplus per traded unit, retail adverse-selection rate, welfare dispersion, and surplus-transfer gap.
- **Safety metrics**: negative-balance violations, conservation breaches, and risk rejections.

7 Experimental Setup

We report two result layers. The first is a single-seed reproducibility snapshot. The second is an eight-seed aggregate profile that is more appropriate for paper-facing discussion. The aggregate profile uses seeds

$$[7, 11, 19, 23, 29, 31, 37, 41].$$

All reported scenarios use the same simulator implementation and differ only in matching mode, batch-window size, seed, and population intensity. The current setup uses a discrete-time step duration of 10 milliseconds.

8 Results

8.1 Aggregate Summary

8.2 Observed Tradeoffs

Four observations are currently the most informative.

First, immediate execution retains the lowest latency profile, but it also carries a high queue-advantage proxy and a large welfare gap. Under the current synthetic environment, low latency alone is not a good indicator of balanced market quality.

Second, the 50 millisecond speed-bump baseline raises latency without reducing queue advantage. It preserves the immediate-style queue signal (0.0742) while worsening the welfare gap to

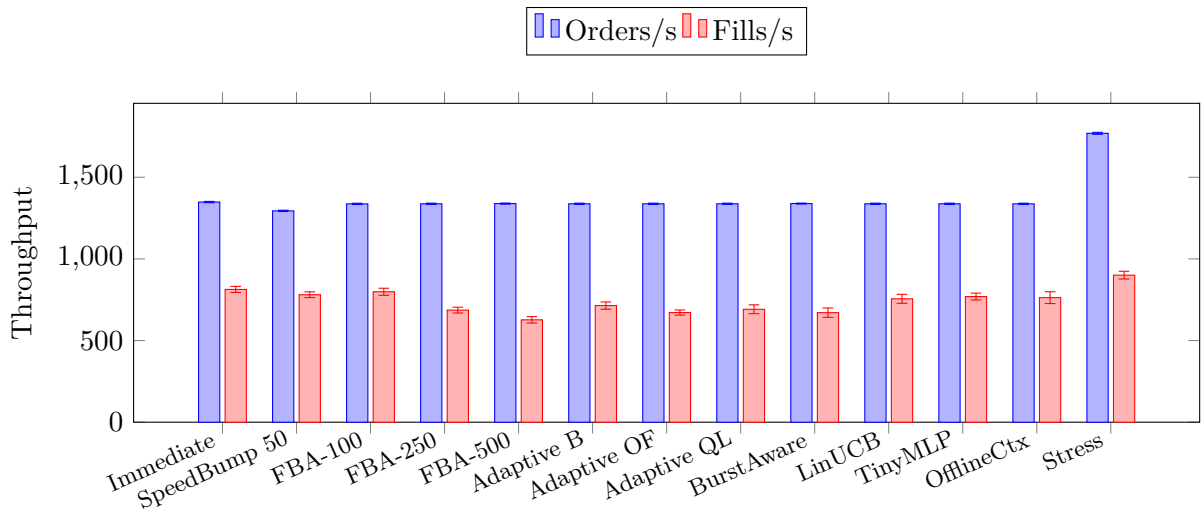

Figure 1: Aggregate throughput comparison with 95% confidence-interval error bars.

Scenario	Orders/s	Fills/s	p50 (ms)	p99 (ms)
Immediate-Surrogate	1348.2	813.1	10.0	10.0
SpeedBump-50ms	1294.3	780.6	60.0	60.0
FBA-100ms	1337.1	798.7	46.3	146.3
FBA-250ms	1337.6	686.4	97.5	452.5
FBA-500ms	1338.8	626.9	213.8	835.0
Adaptive-100-250ms	1337.6	714.3	81.3	360.0
Policy-BurstAware	1338.9	670.8	98.8	400.0
Policy-LearnedLinUCB	1337.6	755.7	48.8	155.0
Policy-LearnedTinyMLP	1337.6	769.3	62.5	221.3
Policy-LearnedOfflineContextual	1337.4	762.8	65.0	215.0
FBA-250ms-Stress	1769.3	900.5	97.5	373.8

Table 1: Eight-seed aggregate summary. All runs report zero conservation breaches and zero negative-balance violations.

4.1034. This makes it a useful negative control: latency rises, but batching-style fairness does not automatically follow.

Third, the fixed 250 millisecond batch regime lowers queue advantage to 0.0273 and compresses the welfare gap relative to immediate execution, but it gives up substantial tail latency. The balanced adaptive heuristic is the strongest fixed-logic mechanism in the current repo on combined welfare terms: it keeps price impact low (4.71), reduces arbitrage-profit proxy to 522.00, and nearly closes the welfare gap (0.0278).

Fourth, the controller comparison now exposes a clearer tradeoff than the earlier draft. `LinUCB` is still best on tail latency (155.0 ms) and strong on fill throughput, but it does so by widening both queue advantage and the surplus-transfer gap. `TinyMLP` improves fills further, but still leaves retail surplus negative and keeps arbitrage capture high. The offline contextual controller is the strongest balanced learned baseline in the current repo: compared with `LinUCB`, it gives up some tail latency but cuts price impact (4.94 versus 5.90), lowers queue advantage (0.0294 versus 0.0513), lowers

Scenario	Impact	Queue Adv.	Retail Surplus	Welfare Gap
Immediate-Surrogate	3.38	0.0742	-0.3710	2.0430
FBA-250ms	5.36	0.0273	-0.3493	0.8896
Adaptive-100-250ms	4.71	0.0375	-0.1508	0.0278
Policy-BurstAware	5.31	0.0305	-0.2258	0.7579
Policy-LearnedLinUCB	5.90	0.0513	0.0795	2.1694
Policy-LearnedTinyMLP	5.22	0.0336	-0.3128	1.9719
Policy-LearnedOfflineContextual	4.94	0.0294	-0.1090	1.3769
FBA-250ms-Stress	5.35	0.0269	-0.1494	0.3805

Table 2: Aggregate market-quality, welfare, and fairness-adjacent summary metrics.

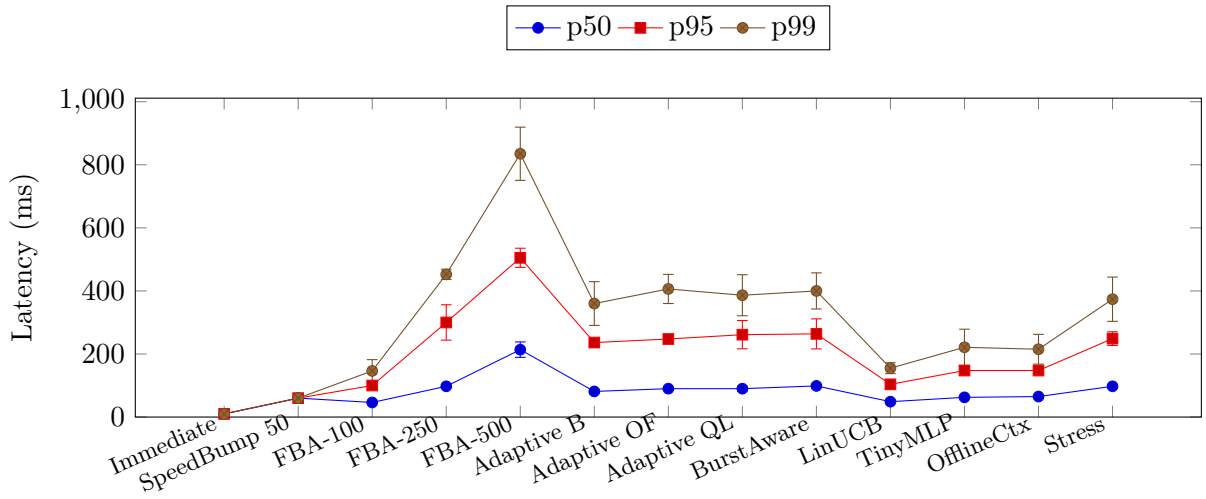

Figure 2: Latency profiles with 95% confidence intervals.

arbitrage-profit proxy (771.25 versus 976.63), and shrinks the welfare gap (1.3769 versus 2.1694).

The stress scenario raises throughput to 1769.25 orders per second and 900.50 fills per second, but it also raises arbitrage capture to 2057.00. Higher activity does not automatically imply better fairness or welfare outcomes.

Across all $13 \times 8 = 104$ aggregate runs, the environment reports zero negative-balance violations and zero conservation breaches.

8.3 Ablation and Sweep Snapshots

We generate a targeted market-structure ablation suite over seeds $[13, 17, 19, 23]$ for a stressed constrained 250 millisecond batch regime. Relaxing risk limits increases throughput from 1190.48 to 1770.24 orders per second and removes all risk rejections, but it also increases the arbitrage-profit proxy from 685.25 to 2198.50. Randomizing tie-breaks leaves throughput roughly unchanged while worsening p99 latency from 320.00 to 402.50 milliseconds. Disabling settlement checks does not change throughput in this setup, indicating that invariant enforcement is not the dominant bottleneck for the present benchmark scale.

We also generate an agent/workload sweep over seeds $[43, 47, 53, 59]$. Removing arbitrageurs collapses arbitrage-profit proxy to zero and flips queue advantage negative. Doubling arbitrage

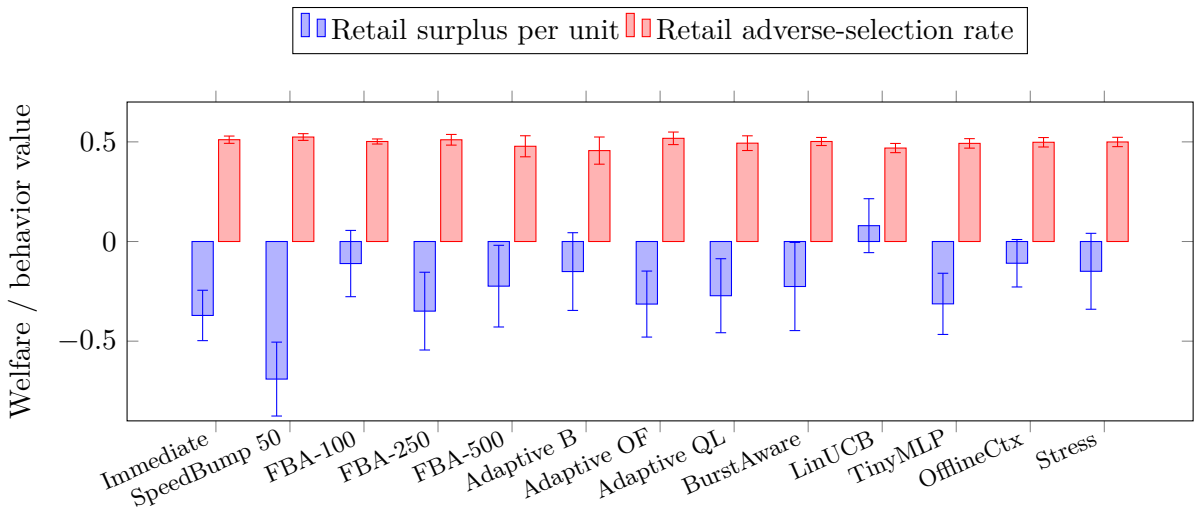

Figure 3: Welfare and behavior signals with 95% confidence intervals. Controllers that improve latency do not automatically improve retail outcomes.

intensity pushes arbitrage-profit proxy to 1730.75. Widening market-maker quotes pushes p99 to 485.00 milliseconds and lowers fills to 595.63. Increasing informed intensity raises throughput modestly while leaving arbitrage-profit much closer to control than the arbitrage sweep.

Finally, we expose a unified $4 \times 3 \times 3 \times 3$ hypercube over arbitrageur intensity $\{0, 1, 2, 3\}$, retail intensity $\{1, 2, 3\}$, informed intensity $\{1, 2, 3\}$, and maker quote width $\{1, 2, 3\}$. Two selected contrasts make the point clearly. At $(\text{arb} = 0, \text{retail} = 1, \text{informed} = 2, \text{maker} = 1)$, the environment reports 1350.60 orders per second, zero arbitrage-profit proxy, and a welfare gap of 0.2376. At $(\text{arb} = 3, \text{retail} = 1, \text{informed} = 2, \text{maker} = 1)$, throughput rises to 1542.86 and p99 falls to 362.50 milliseconds, but arbitrage-profit proxy jumps to 1982.00 and the welfare gap widens to 1.2839. The hypercube therefore makes it possible to separate "high activity" from "high transfer-to-arbitrageur" regimes under one artifact family.

9 Limitations

The current artifact is stronger than the earlier benchmark scaffold, but it is still short of a mature top-tier benchmark submission. The most important limitations are:

- the learned-controller family is still lightweight and discrete-action
- welfare metrics are derived from a synthetic fundamental rather than real market replay
- the hypercube is slice-visualized rather than summarized by a more compact statistical surface
- the agent behaviors are stylized and do not yet include richer strategic adaptation

10 Related Work

The benchmark is motivated directly by frequent-batch-auction market design [2]. Its benchmark-oriented framing is also aligned with the broader machine-learning practice of exposing reusable environments rather than only single-task scripts [1]. The present contribution sits between those

two traditions: it is neither a pure market-design theory paper nor a standard reinforcement-learning benchmark, but an infrastructure-aware environment that aims to make later learning or control work more credible.

11 Conclusion

This NeurIPS-track benchmark line now supports a more defensible benchmark narrative than the earlier scaffold. It preserves settlement invariants, exposes a reusable interaction API, adds a stronger offline contextual baseline, expands fairness measurement into welfare and behavior signals, and merges arbitrage pressure into a unified higher-dimensional stress surface. The current results show a clear tradeoff: controllers that optimize aggressively for latency and fills widen the surplus-transfer gap, while more balanced controllers can give up some tail performance to materially improve market-quality and welfare outcomes.

A Appendix Tables

Scenario	Orders/s	p99 (ms)	Arb Profit	Risk Rejects
Ablation-Control	1190.48	320.00	685.25	2922
Ablation-RelaxedRisk	1770.24	302.50	2198.50	0
Ablation-RandomTieBreak	1190.48	402.50	738.50	2908
Ablation-NoSettlementChecks	1190.48	320.00	685.25	2922

Table 3: Mechanism ablation appendix table.

Controller	p99 (ms)	Impact	Retail Surplus	Welfare Gap
BurstAware	400.00	5.31	-0.2258	0.7579
LinUCB	155.00	5.90	0.0795	2.1694
TinyMLP	221.25	5.22	-0.3128	1.9719
OfflineContextual	215.00	4.94	-0.1090	1.3769

Table 4: Controller comparison appendix table.

Arb	Retail	Informed	Maker	Welfare Gap
0	1	2	1	0.2376
3	1	2	1	1.2839
0	3	2	1	0.1335
3	3	2	3	1.8005

Table 5: Selected cells from the unified arbitrage / retail / informed / maker hypercube.

References

[1] Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. Openai gym, 2016.

[2] Eric Budish, Peter Cramton, and John Shim. The high-frequency trading arms race: Frequent batch auctions as a market design response. *The Quarterly Journal of Economics*, 130(4):1547–1621, 2015.