A Ledger-Aware Benchmark for Market Infrastructure under Frequent Batch Auctions

Author Name
Institution
email@example.com

March 8, 2026

Abstract

We present a ledger-aware benchmark track for studying how market mechanisms and lightweight learned controllers trade off latency, fill throughput, and retail outcomes under explicit settlement constraints. We explicitly study how learning-based controllers operate under settlement-constrained market environments rather than in a matching-only simulator. The environment combines seedable agent-based order flow, configurable immediate, speed-bump, and frequent-batch-auction matching regimes, double-entry style account transitions, and deterministic safety checks for conservation and non-negativity. The paper-facing welfare decomposition is intentionally narrow: retail surplus per traded unit, retail adverse-selection rate, and surplus-transfer gap are treated as the primary outcome family, while broader diagnostics remain in the artifact layer. The current version exposes a step-wise `Reset/Step/Observe/Metrics` API, an adapter-driven control surface, five policy baselines including offline contextual and fitted-Q controllers, and a unified four-dimensional stress sweep over arbitrage intensity, retail intensity, informed intensity, and maker quote width. It also adds a minimal offline-training story with held-out regime evaluation over unseen stress combinations. Across all published runs, settlement invariants remain intact while mechanism and controller choices induce clear latency, fairness, and welfare-transfer tradeoffs.

1 Introduction

Electronic markets are shaped jointly by matching rules, latency structure, and settlement semantics. Market-design work on frequent batch auctions argues that short batching windows can reduce some latency-driven distortions in continuous markets [2]. In practice, many evaluation artifacts study matching rules in isolation and do not model the settlement layer, replay behavior, or safety checks that an actual market stack must satisfy.

This paper develops a benchmark-oriented layer on top of an existing ledger-first market-infrastructure prototype. It is not a replacement for the original systems paper. It is a separate evaluation track with a narrower question: can market-design and controller tradeoffs be measured inside an environment that preserves explicit settlement invariants and summarizes retail outcomes through a small, interpretable welfare decomposition? In particular, we study how learning-based controllers behave when settlement invariants, replay stability, and market-design choices are all part of the environment definition.

The current paper makes four concrete contributions:

1. It defines a ledger-aware benchmark environment in which market-design experiments are coupled to deterministic settlement checks.

2. It provides a seedable multi-agent order-flow generator, a unified four-dimensional stress surface over arbitrage, retail, informed, and maker-width multipliers, and a compact hypercube summary over the same surface.

3. It reports single-seed and eight-seed aggregate results over latency/fill efficiency and a paper-facing welfare decomposition: retail surplus, retail adverse selection, and surplus-transfer gap.

4. It preserves executable settlement invariants across all artifacts, so mechanism and controller improvements are evaluated under non-negativity and conservation constraints rather than in a matching-only simulator.

2 Problem Setting

We study a discrete-time environment in which a market state is advanced one simulator step at a time. At step t, the state is represented as

$$s_t = (B_t, S_t, X_t, f_t),$$

where B_t and S_t denote the buy and sell books, X_t is the vector of account states, and f_t is the latent synthetic fundamental used by the order-flow generator. A scenario additionally fixes a matching mode, a batch-window size, risk thresholds, and a population of agents.

The benchmark focuses on four concrete questions:

1. How do immediate, speed-bump, adaptive, and frequent-batch-auction execution differ in latency and fill behavior?

2. Can fairness-adjacent signals such as queue-priority advantage and latency-arbitrage profit be measured reproducibly under deterministic settlement semantics?

3. Do welfare/behavior metrics such as retail surplus and adverse selection reveal controller trade-offs that proxy metrics alone miss?

4. Do settlement invariants hold under heterogeneous agents and stress scenarios?

3 Benchmark Environment

3.1 Core Components

The benchmark environment lives in the repository's `simulator/` package. It consists of scenario, order, fill, and result schemas, heterogeneous agent order generation, immediate, speed-bump, adaptive, and batch-clearing matching paths, deterministic settlement application, metric accumulation, and benchmark artifact writers.

The environment also exposes a step-wise control surface with `Reset()`, `Observe()`, `Step()`, and `Metrics()`. On top of that surface, the repository includes an adapter layer that returns reward-bearing timesteps and advertises per-scenario action support. The current implementation exposes five runtime control channels: batch-window override for adaptive scenarios, risk-limit scaling, tie-break randomization, release-cadence control, and price-aggression bias [1].

3.2 Determinism and Invariants

Each scenario is seeded. Given the same scenario configuration and seed, the environment is expected to produce the same accepted order counts, fill counts, and benchmark metrics. The benchmark currently enforces three operational invariants:

1. **Conservation**: total cash and total inventory across accounts remain constant over time.
2. **Non-negativity**: no account is allowed to end a fill application with negative cash or negative units.
3. **Deterministic replay**: seeded runs produce stable outputs for the same configuration.

4 Agent Models

The present environment includes four agent classes:

- **Market makers**: quote on both sides around the current fundamental with configurable width and intensity.
- **Latency arbitrageurs**: react to stale price information and attempt to exploit short-horizon dislocations.
- **Retail traders**: submit noisier buy or sell orders with wider dispersion around the current reference level.
- **Informed traders**: trade in the direction of the next-step fundamental when the generator exposes a directional signal.

These agents are intentionally simple. Their purpose is to generate heterogeneous but controlled activity so that matching rules and controller choices can be compared under a repeatable workload.

5 Tasks and Baselines

The current benchmark suite contains fourteen scenarios:

1. `Immediate-Surrogate`
2. `SpeedBump-50ms`
3. `FBA-100ms`
4. `FBA-250ms`
5. `FBA-500ms`
6. `Adaptive-100-250ms`
7. `Adaptive-OrderFlow-100-250ms`
8. `Adaptive-QueueLoad-100-250ms`
9. `Policy-BurstAware-100-250ms`

10. `Policy-LearnedLinUCB-100-250ms`
11. `Policy-LearnedTinyMLP-100-250ms`
12. `Policy-LearnedOfflineContextual-100-250ms`
13. `Policy-LearnedFittedQ-100-250ms`
14. `FBA-250ms-Stress`

The controller suite now contains a hand-written burst-aware controller, a contextual linear bandit, a gradient-trained TinyMLP policy, an offline contextual value controller, and an offline fitted-Q controller. The offline controllers are trained from logged trajectories generated by burst-aware, LinUCB, TinyMLP, offline contextual, and random behavior policies.

6 Metrics

The environment reports five groups of signals:

- **Systems metrics**: orders per second, fills per second, and p50/p95/p99 latency.
- **Market-quality metrics**: average spread and average price impact.
- **Fairness-adjacent proxies**: queue-priority advantage, latency-arbitrage profit, and execution dispersion.
- **Primary welfare decomposition**: retail surplus per traded unit, retail adverse-selection rate, and surplus-transfer gap.
- **Secondary diagnostics**: arbitrageur surplus per traded unit and welfare dispersion.
- **Safety metrics**: negative-balance violations, conservation breaches, and risk rejections.

7 Experimental Setup

We report two result layers. The first is a single-seed reproducibility snapshot. The second is an eight-seed aggregate profile that is more appropriate for paper-facing discussion. The aggregate profile uses seeds

$$[7, 11, 19, 23, 29, 31, 37, 41].$$

The learned fitted-Q controller uses an explicit offline train/eval split. Logged trajectories are collected on training seeds $[181, 191, 193, 197, 199, 211]$ from burst-aware, LinUCB, TinyMLP, offline contextual, and random behavior policies. The fitted-Q loop runs for eight iterations with discount 0.97 over the shared discrete action bundle. Held-out evaluation uses seeds $[223, 227, 229, 233]$ on four unseen stress regimes defined in the artifact layer. All reported scenarios use the same simulator implementation and differ only in matching mode, batch-window size, seed, and population intensity. The current setup uses a discrete-time step duration of 10 milliseconds.

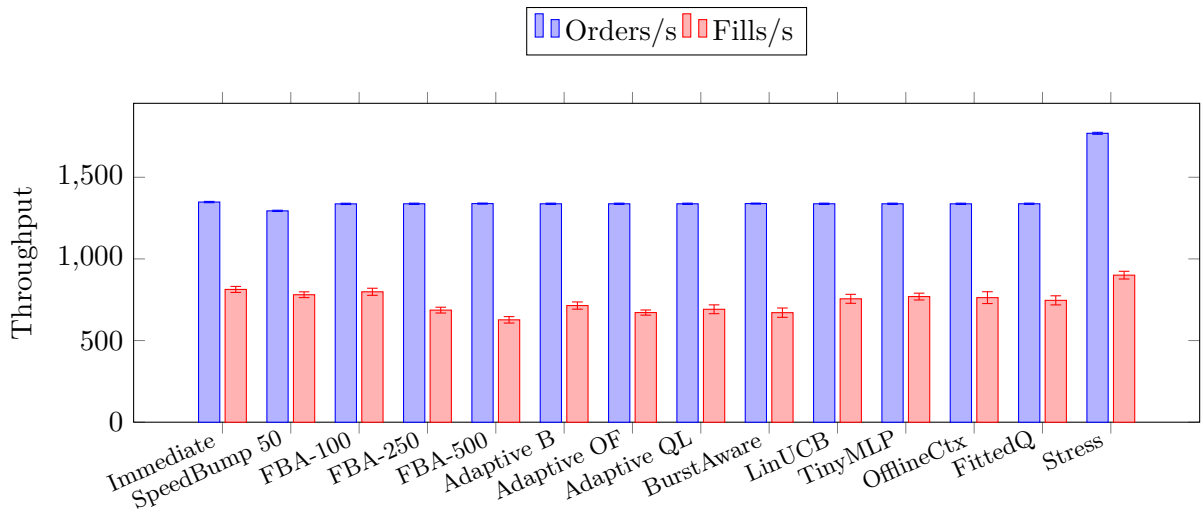

Figure 1: Aggregate throughput comparison with 95% confidence-interval error bars.

8 Results

8.1 Aggregate Summary

8.2 Observed Tradeoffs

Four observations are currently the most informative.

First, immediate execution retains the lowest latency profile, but it also carries a high queue-advantage proxy and a large welfare gap. Under the current synthetic environment, low latency alone is not a good indicator of balanced market quality.

Second, the 50 millisecond speed-bump baseline raises latency without reducing queue advantage. It preserves the immediate-style queue signal (0.0742) while worsening the welfare gap to 4.1034. This makes it a useful negative control: latency rises, but batching-style fairness does not automatically follow.

Third, the fixed 250 millisecond batch regime lowers queue advantage to 0.0273 and compresses the welfare gap relative to immediate execution, but it gives up substantial tail latency. The balanced adaptive heuristic is the strongest fixed-logic mechanism in the current repo on combined welfare terms: it keeps price impact low (4.71), reduces arbitrage-profit proxy to 522.00, and nearly closes the welfare gap (0.0278).

Fourth, the controller comparison now exposes a clearer tradeoff than the earlier draft. `LinUCB` remains strong on tail latency and fill throughput, but it does so by widening both queue advantage and the surplus-transfer gap. `FittedQ` now produces the best in-distribution p99 tail (145.0 ms), giving a minimal offline-RL style training story on top of the adapter API, but it still behaves closer to `LinUCB` than to `OfflineContextual` on welfare transfer. `TinyMLP` improves fills further, but still leaves retail surplus negative and keeps arbitrage capture high. The offline contextual controller remains the strongest balanced learned baseline in the current repo: compared with `LinUCB`, it gives up some tail latency but cuts price impact (4.94 versus 5.90), lowers queue advantage (0.0294 versus 0.0513), lowers arbitrage-profit proxy (771.25 versus 976.63), and shrinks the welfare gap (1.3769 versus 2.1694).

The stress scenario raises throughput to 1769.25 orders per second and 900.50 fills per second, but it also raises arbitrage capture to 2057.00. Higher activity does not automatically imply better

Scenario	Orders/s	Fills/s	p50 (ms)	p99 (ms)
Immediate-Surrogate	1348.2	813.1	10.0	10.0
SpeedBump-50ms	1294.3	780.6	60.0	60.0
FBA-100ms	1337.1	798.7	46.3	146.3
FBA-250ms	1337.6	686.4	97.5	452.5
FBA-500ms	1338.8	626.9	213.8	835.0
Adaptive-100-250ms	1337.6	714.3	81.3	360.0
Policy-BurstAware	1338.9	670.8	98.8	400.0
Policy-LearnedLinUCB	1337.6	755.7	48.8	155.0
Policy-LearnedTinyMLP	1337.6	769.3	62.5	221.3
Policy-LearnedOfflineContextual	1337.4	762.8	65.0	215.0
Policy-LearnedFittedQ	1337.7	746.2	46.3	145.0
FBA-250ms-Stress	1769.3	900.5	97.5	373.8

Table 1: Eight-seed aggregate summary. All runs report zero conservation breaches and zero negative-balance violations.

Scenario	Impact	Queue Adv.	Retail Surplus	Welfare Gap
Immediate-Surrogate	3.38	0.0742	-0.3710	2.0430
FBA-250ms	5.36	0.0273	-0.3493	0.8896
Adaptive-100-250ms	4.71	0.0375	-0.1508	0.0278
Policy-BurstAware	5.31	0.0305	-0.2258	0.7579
Policy-LearnedLinUCB	5.90	0.0513	0.0795	2.1694
Policy-LearnedTinyMLP	5.22	0.0336	-0.3128	1.9719
Policy-LearnedOfflineContextual	4.94	0.0294	-0.1090	1.3769
Policy-LearnedFittedQ	5.88	0.0451	0.0742	2.2036
FBA-250ms-Stress	5.35	0.0269	-0.1494	0.3805

Table 2: Aggregate market-quality, welfare, and fairness-adjacent summary metrics.

fairness or welfare outcomes.

Across all $14 \times 8 = 112$ aggregate runs, the environment reports zero negative-balance violations and zero conservation breaches.

8.3 Held-Out Regime Generalization

Held-out results use unseen seeds $[223, 227, 229, 233]$ and four unseen regime combinations. The policy summary is shown in Table 3.

The held-out picture is materially different from the in-distribution aggregate. `FittedQ` lowers both p99 and welfare gap relative to `LinUCB` in the high-arbitrage wide-maker, retail-burst, and composite-stress held-out regimes. In the informed-wide regime, it gives up some tail latency (182.5 versus 167.5 ms) but still lowers welfare gap (2.2994 versus 2.3765). `OfflineContextual` remains the most welfare-balanced learned baseline overall, but it is materially slower than both `LinUCB` and `FittedQ` on held-out tails.

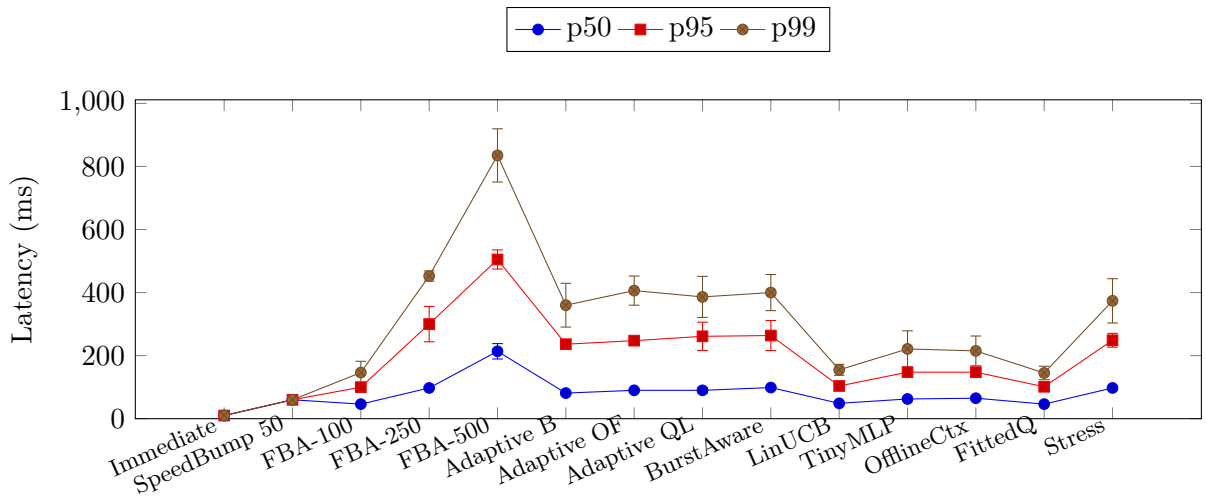

Figure 2: Latency profiles with 95% confidence intervals.

Policy	Fills/s	p99 (ms)	Retail Surplus	Welfare Gap
BurstAware	834.08	376.88	-0.4042	1.6286
LinUCB	978.52	171.25	0.1215	2.4466
OfflineContextual	961.61	275.62	-0.0732	1.9535
FittedQ	946.78	159.38	0.1110	2.3158

Table 3: Held-out policy summary over four unseen regimes.

8.4 Fitted-Q Learning Curve

The fitted-Q controller is no longer just a point baseline. We also log its held-out behavior over Bellman-update iterations. The untrained snapshot begins at a welfare gap of 3.2117 ± 0.5538 and p99 latency of 341.25 ± 22.97 ms. After the first Bellman update, the held-out welfare gap falls to 2.0790 ± 0.4767 and p99 falls to 198.75 ± 25.39 ms. By iteration 8, Bellman MSE declines from 45.1571 to 6.5753, held-out p99 reaches 155.62 ± 12.84 ms, and the welfare gap settles at 2.4226 ± 0.5400. The training story is therefore not monotone on every objective: later iterations improve Bellman error and tail latency, but they give back part of the earliest welfare improvement.

8.5 Ablation and Sweep Snapshots

We generate a targeted market-structure ablation suite over seeds $[13, 17, 19, 23]$ for a stressed constrained 250 millisecond batch regime. Relaxing risk limits increases throughput from 1190.48 to 1770.24 orders per second and removes all risk rejections, but it also increases the arbitrage-profit proxy from 685.25 to 2198.50. Randomizing tie-breaks leaves throughput roughly unchanged while worsening p99 latency from 320.00 to 402.50 milliseconds. Disabling settlement checks does not change throughput in this setup, indicating that invariant enforcement is not the dominant bottleneck for the present benchmark scale.

We also generate an agent/workload sweep over seeds $[43, 47, 53, 59]$. Removing arbitrageurs collapses arbitrage-profit proxy to zero and flips queue advantage negative. Doubling arbitrage intensity pushes arbitrage-profit proxy to 1730.75. Widening market-maker quotes pushes p99 to 485.00 milliseconds and lowers fills to 595.63. Increasing informed intensity raises throughput

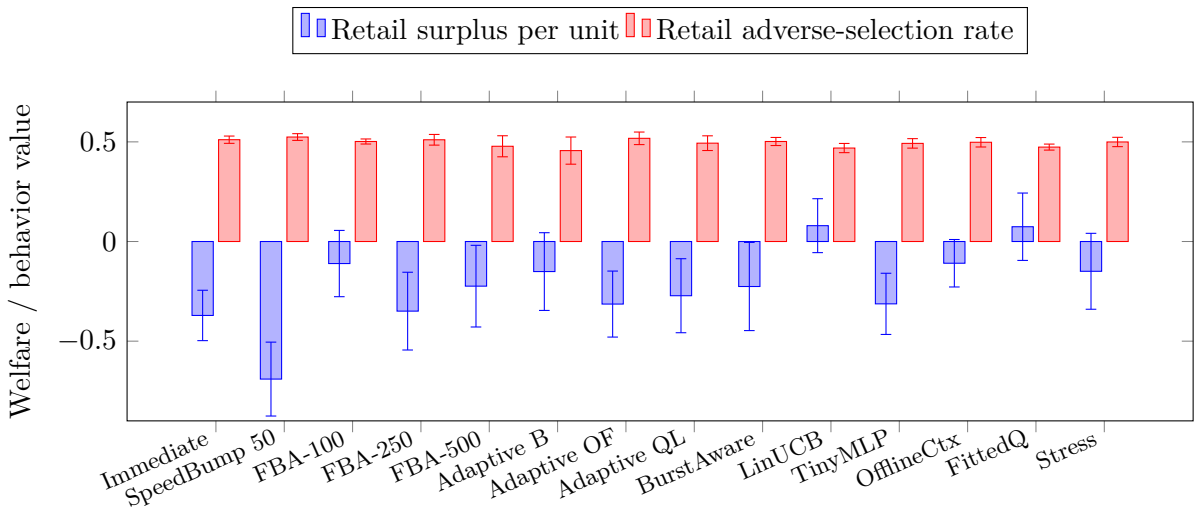

Figure 3: Welfare and behavior signals with 95% confidence intervals. Controllers that improve latency do not automatically improve retail outcomes.

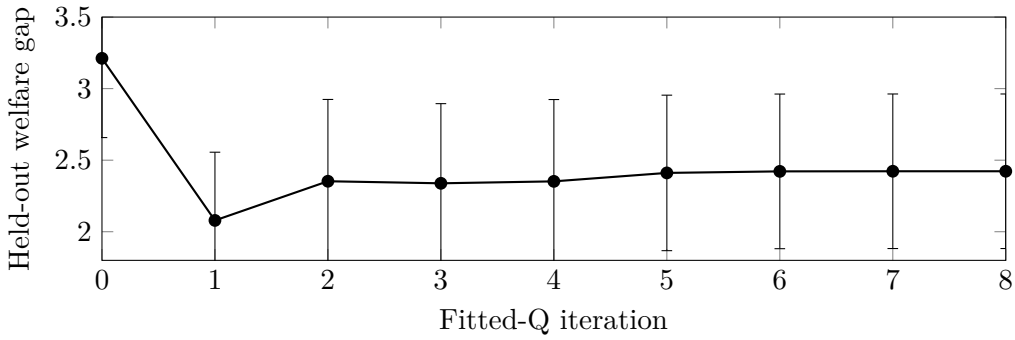

Figure 4: Held-out fitted-Q learning curve. The first Bellman update captures most of the welfare-gap gain; later iterations continue to reduce Bellman error and tail latency but do not preserve that full early welfare improvement.

modestly while leaving arbitrage-profit much closer to control than the arbitrage sweep.

Finally, we expose a unified $4 \times 3 \times 3 \times 3$ hypercube over arbitrageur intensity $\{0, 1, 2, 3\}$, retail intensity $\{1, 2, 3\}$, informed intensity $\{1, 2, 3\}$, and maker quote width $\{1, 2, 3\}$. Two raw cells still illustrate the environment directly. At $(\text{arb} = 0, \text{retail} = 1, \text{informed} = 2, \text{maker} = 1)$, the environment reports 1350.60 orders per second, zero arbitrage-profit proxy, and a welfare gap of 0.2376. At $(\text{arb} = 3, \text{retail} = 1, \text{informed} = 2, \text{maker} = 1)$, throughput rises to 1542.86 and p99 falls to 362.50 milliseconds, but arbitrage-profit proxy jumps to 1982.00 and the welfare gap widens to 1.2839.

This compact summary removes the need to interpret the hypercube only through slice families. Arbitrage pressure is the dominant driver of welfare-gap expansion, retail intensity is mostly a throughput lever, and wider maker quotes worsen retail outcome without moving aggregate activity in this setup. Averaged over informed intensity and maker width, moving from arb = 0 to arb = 3 widens the welfare gap by 1.1748, 1.2262, and 1.2287 at retail intensities $x1$, $x2$, and $x3$, respectively. This effect persists across retail intensities. Higher retail flow does not neutralize transfer-to-

Factor	Δ Orders/s	Δ Retail Surplus	Δ Retail Adverse	Δ Welfare Gap
Arbitrage $0 \to 3$	176.26	-0.1804	0.0117	1.2099
Retail $1 \to 3$	780.94	0.0772	0.0018	0.0781
Informed $1 \to 3$	150.91	-0.1986	0.0131	0.0876
Maker width $1 \to 3$	0.00	-0.1250	0.0085	0.2653

Table 4: Compact high-low contrasts from the unified hypercube.

arbitrageur.

9 Limitations

The current artifact is stronger than the earlier benchmark scaffold, but it is still short of a mature top-tier benchmark submission. The most important limitations are:

- the learned-controller family is still lightweight and discrete-action
- welfare metrics are derived from a synthetic fundamental rather than real market replay
- the hypercube is summarized by low-order main effects and contrasts rather than by a fully fitted statistical model
- the agent behaviors are stylized and do not yet include richer strategic adaptation

10 Related Work

The benchmark is motivated directly by frequent-batch-auction market design [2]. Its benchmark-oriented framing is also aligned with the broader machine-learning practice of exposing reusable environments rather than only single-task scripts [1]. The present contribution sits between those two traditions: it is neither a pure market-design theory paper nor a standard reinforcement-learning benchmark, but an infrastructure-aware environment that aims to make later learning or control work more credible.

11 Conclusion

This NeurIPS-track benchmark line now supports a tighter benchmark narrative than the earlier scaffold. It preserves settlement invariants, exposes a reusable interaction API, reduces the paper-facing welfare story to a small interpretable decomposition, and compresses the unified stress surface into main effects and contrasts that are easier to reason about. The current results show a clear tradeoff: controllers that optimize aggressively for latency and fills widen the surplus-transfer gap, while more balanced controllers can give up some tail performance to materially improve market-quality and retail outcomes.

Scenario	Orders/s	p99 (ms)	Arb Profit	Risk Rejects
Ablation-Control	1190.48	320.00	685.25	2922
Ablation-RelaxedRisk	1770.24	302.50	2198.50	0
Ablation-RandomTieBreak	1190.48	402.50	738.50	2908
Ablation-NoSettlementChecks	1190.48	320.00	685.25	2922

Table 5: Mechanism ablation appendix table.

Controller	p99 (ms)	Impact	Retail Surplus	Welfare Gap
BurstAware	400.00	5.31	-0.2258	0.7579
LinUCB	155.00	5.90	0.0795	2.1694
TinyMLP	221.25	5.22	-0.3128	1.9719
OfflineContextual	215.00	4.94	-0.1090	1.3769
FittedQ	145.00	5.88	0.0742	2.2036

Table 6: Controller comparison appendix table.

A Appendix Tables

References

[1] Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. Openai gym, 2016.

[2] Eric Budish, Peter Cramton, and John Shim. The high-frequency trading arms race: Frequent batch auctions as a market design response. *The Quarterly Journal of Economics*, 130(4):1547–1621, 2015.

Arb	Retail	Informed	Maker	Welfare Gap
0	1	2	1	0.2376
3	1	2	1	1.2839
0	3	2	1	0.1335
3	3	2	3	1.8005

Table 7: Selected cells from the unified arbitrage / retail / informed / maker hypercube.