A Ledger-Aware Benchmark for Market Infrastructure under Frequent Batch Auctions

Author Name
Institution
email@example.com

March 7, 2026

Abstract

We present a benchmark-oriented extension of a ledger-first market infrastructure prototype for studying market-design and execution behavior under settlement and risk constraints. The environment combines seedable agent-based order flow, configurable immediate, speed-bump, and frequent-batch-auction matching regimes, double-entry style account state transitions, and deterministic safety checks for conservation and non-negativity. We expose reproducible benchmark tasks that compare latency, throughput, spread, price impact, queue-priority advantage, and latency-arbitrage profit across market designs. In the current implementation, all generated scenarios preserve settlement invariants, while mechanism changes induce clear latency and market-quality tradeoffs. The result is a benchmark draft for market-infrastructure research that is separate from the original systems-paper line and is intended as a foundation for stronger benchmark-style evaluation.

1 Introduction

Electronic markets are shaped jointly by matching rules, latency structure, and settlement semantics. Market-design work on frequent batch auctions argues that short batching windows can reduce some latency-driven distortions in continuous markets [2]. In practice, however, many evaluation artifacts study matching rules in isolation and do not model the settlement layer, replay behavior, or safety checks that an actual market stack must satisfy. That omission matters because execution quality can improve in a simulator while the corresponding state transitions still violate conservation or allow invalid account evolution.

This paper develops a benchmark-oriented layer on top of an existing ledger-first market-infrastructure prototype. The benchmark is not a replacement for the original systems paper. Instead, it is a separate evaluation track that asks a narrower question: can we measure market-design tradeoffs inside an environment that preserves explicit settlement invariants? Our current artifact implements seedable heterogeneous agents, immediate, speed-bump, adaptive, and fixed-window batch matching, deterministic account updates, and generated benchmark outputs that can be reproduced from the repository.

The benchmark is motivated by a research gap between market-design arguments and reproducible infrastructure-aware evaluation. It is also informed by the broader benchmark culture in machine learning, where a reusable environment and a clear task interface often matter as much as a single model contribution [1]. Our present contribution is not a new learning algorithm. It is a benchmark scaffold that exposes the right constraints for later policy-learning, control, or mechanism-comparison work.

The current paper makes four concrete contributions:

1. It defines a ledger-aware benchmark environment in which market-design experiments are coupled to deterministic settlement checks.
2. It provides a seedable multi-agent order-flow generator with four agent classes and ten benchmark scenarios.
3. It reports single-seed and eight-seed aggregate benchmark results over throughput, latency, spread, price impact, queue-priority advantage, and arbitrage-profit proxies.
4. It exposes a step-wise `Reset/Step/Observe/Metrics` API, a minimal gym-style adapter with batch-window, risk-scale, and tie-break controls, an adapter-driven policy baseline, and a first ablation suite for risk limits, tie-breaking, and settlement checks.

2 Problem Setting

We study a discrete-time environment in which a market state is advanced one simulator step at a time. At step t, the state is represented as

$$s_t = (B_t, S_t, X_t, f_t),$$

where B_t and S_t denote the current buy and sell books, X_t is the vector of account states, and f_t is the latent fundamental used by the synthetic order-flow generator. A scenario additionally fixes a matching mode, a batch-window size, risk thresholds, and a population of agents.

The benchmark focuses on the following research questions:

1. How do immediate, speed-bump, adaptive, and frequent-batch-auction execution differ in latency and fill behavior?
2. Can fairness-adjacent signals such as queue-priority advantage and latency-arbitrage profit be measured reproducibly under deterministic settlement semantics?
3. Do settlement invariants hold under heterogeneous agents and stress scenarios?

The benchmark output is not a reward signal for a trained policy. Instead, it is a structured set of metrics that later work can use either as direct benchmark targets or as components in a downstream objective. This is deliberate: the current stage of the project aims to establish a stable evaluation layer before introducing adaptive control or learned market policies.

3 Benchmark Environment

3.1 Core Components

The benchmark environment lives in the repository's `simulator/` package. It consists of:

- scenario, order, fill, and result schemas
- heterogeneous agent order generation
- immediate, speed-bump, adaptive, and batch-clearing matching paths
- deterministic settlement application

- metric accumulation and artifact generation

The environment is designed so that matching and settlement are not independent stages. Every fill is immediately translated into account updates, and every run is checked for safety conditions before benchmark artifacts are accepted. This design choice is the main difference between the present artifact and a lightweight matching-only simulator.

The current environment also exposes a step-wise control surface with `Reset()`, `Observe()`, `Step()`, and `Metrics()`, which allows later work to treat the benchmark as a control loop rather than only a batch artifact. On top of that surface, the repository now includes a minimal adapter layer that returns reward-bearing timesteps and advertises per-scenario action support. In the current implementation, the adapter exposes three explicit control channels: batch-window override for adaptive scenarios, risk-limit scaling, and tie-break randomization for batch-style matching.

3.2 Determinism and Invariants

Each scenario is seeded. Given the same scenario configuration and seed, the environment is expected to produce the same accepted order counts, fill counts, and proxy metrics. The test suite includes a determinism regression for this requirement.

The benchmark currently enforces three operational invariants:

1. **Conservation**: total cash and total inventory across accounts remain constant over time.
2. **Non-negativity**: no account is allowed to end a fill application with negative cash or negative units.
3. **Deterministic replay**: seeded runs produce stable outputs for the same configuration.

These are not theorem-level guarantees for every possible extension of the environment. They are executable safety properties of the current artifact, and they are checked as part of benchmark generation.

4 Agent Models

The present environment includes four agent classes:

- **Market makers**: quote on both sides around the current fundamental with configurable width and intensity.
- **Latency arbitrageurs**: react to stale price information and attempt to exploit short-horizon dislocations.
- **Retail traders**: submit noisier buy or sell orders with wider dispersion around the current reference level.
- **Informed traders**: trade in the direction of the next-step fundamental when the generator exposes a directional signal.

These agents are intentionally simple. They are not meant to be realistic economic models of all order flow. Their purpose is to generate heterogeneous but controlled activity so that matching rules can be compared under a repeatable workload. The stress scenario extends the default population with additional retail and arbitrage activity to increase order intensity.

5 Tasks and Baselines

The current benchmark suite contains ten scenarios:

1. `Immediate-Surrogate`
2. `SpeedBump-50ms`
3. `FBA-100ms`
4. `FBA-250ms`
5. `FBA-500ms`
6. `Adaptive-100-250ms`
7. `Adaptive-OrderFlow-100-250ms`
8. `Adaptive-QueueLoad-100-250ms`
9. `Policy-BurstAware-100-250ms`
10. `FBA-250ms-Stress`

The first scenario acts as a continuous-style surrogate. The second applies a fixed 50 millisecond speed bump before immediate matching. The next three fix batch windows at 100, 250, and 500 milliseconds. The sixth scenario applies a balanced heuristic controller that adapts between 100 and 250 milliseconds based on accepted flow and queue depth. The next two scenarios isolate order-flow-driven and queue-load-aware adaptive policies over the same window range. The ninth scenario is an adapter-driven burst-aware controller that acts through explicit control actions rather than only built-in matching logic. The final scenario keeps the 250 millisecond window while increasing the order-flow intensity. These scenarios are modest compared with a mature benchmark suite, but they are enough to expose first-order tradeoffs between latency, fill throughput, and fairness-adjacent proxy metrics.

The benchmark still lacks one baseline that would materially strengthen the NeurIPS-track story:

- a learned or policy-optimized adaptive controller

That is the most important next addition for future work.

6 Metrics

The environment reports four groups of signals.

6.1 Systems Metrics

These metrics capture operational behavior:

- orders per second
- fills per second
- p50, p95, and p99 latency

6.2 Market-Quality Metrics

The current artifact also measures:

- average spread
- average price impact relative to the latent fundamental

6.3 Fairness-Adjacent Proxies

The benchmark is not yet a full fairness-study environment, but it does expose:

- queue-priority advantage
- latency-arbitrage profit
- execution dispersion across agent classes

6.4 Safety Metrics

Finally, the environment reports explicit safety failures:

- negative-balance violations
- conservation breaches
- risk rejections

7 Experimental Setup

We report two result layers. The first is a single-seed reproducibility snapshot. The second is an eight-seed aggregate profile that is more appropriate for paper-facing discussion. The aggregate profile uses seeds

$$[7, 11, 19, 23, 29, 31, 37, 41].$$

All reported scenarios use the same simulator implementation and differ only in matching mode, batch-window size, seed, and population intensity. The current setup uses a discrete-time step duration of 10 milliseconds. Fixed-window batch scenarios aggregate accepted orders until the configured window boundary and then clear them in batch. The immediate surrogate processes each accepted order against the current book as soon as it arrives.

8 Results

8.1 Aggregate Summary

8.2 Observed Tradeoffs

Three measured observations are currently the most informative.

First, immediate execution retains the lowest latency profile, with mean p50, p95, and p99 all equal to 10 milliseconds under the present discrete-time environment. This is expected because the immediate surrogate does not accumulate orders across a multi-step batch horizon.

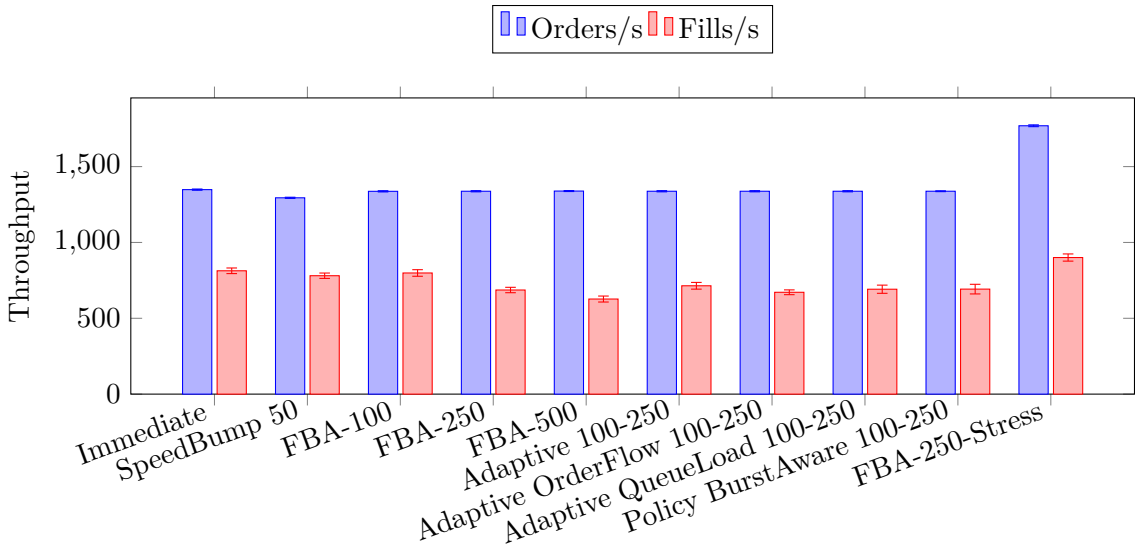

Figure 1: Aggregate throughput comparison with 95% confidence-interval error bars.

Scenario	Orders/s	Fills/s	p50 (ms)	p99 (ms)
Immediate-Surrogate	1348.2	813.1	10.0	10.0
SpeedBump-50ms	1294.3	780.6	60.0	60.0
FBA-100ms	1337.1	798.7	46.3	146.3
FBA-250ms	1337.6	686.4	97.5	452.5
FBA-500ms	1338.8	626.9	213.8	835.0
Adaptive-100-250ms	1337.6	714.3	81.3	360.0
Adaptive-OrderFlow-100-250ms	1337.6	671.4	90.0	406.3
Adaptive-QueueLoad-100-250ms	1337.6	691.6	90.0	386.3
Policy-BurstAware-100-250ms	1337.7	692.3	98.8	406.3
FBA-250ms-Stress	1769.3	900.5	97.5	373.8

Table 1: Eight-seed aggregate summary. All runs report zero conservation breaches and zero negative-balance violations.

Second, the 50 millisecond speed-bump baseline introduces a fixed 60 millisecond latency profile and lowers throughput to 1294.30 orders per second while preserving the same queue-priority-advantage proxy observed under immediate execution. This makes it a useful intermediate baseline: latency rises, but batching-style fairness proxies do not automatically improve.

Third, moving from immediate execution to fixed batch windows creates a clear latency tradeoff without materially changing mean order throughput in the non-stress scenarios. The 100 millisecond batch already raises mean p99 latency to 146.25 milliseconds, while the 500 millisecond batch pushes mean p99 latency to 835.00 milliseconds.

Fourth, the 250 millisecond batch regime lowers the queue-priority-advantage proxy to 0.0273, below immediate execution (0.0742), the speed-bump baseline (0.0742), and the 100 millisecond batch regime (0.0571). The balanced adaptive heuristic settles around a mean window of 207.14 milliseconds and lowers the arbitrage-profit proxy to 522.00, below both fixed 100 and fixed 250 millisecond batching. The order-flow adaptive variant pushes to a higher mean window of 216.67

Scenario	Spread	Impact	Queue Adv.	Arb Profit
Immediate-Surrogate	1.98	3.38	0.0742	669.62
SpeedBump-50ms	1.96	5.96	0.0742	1346.25
FBA-100ms	1.00	5.96	0.0571	1015.75
FBA-250ms	1.00	5.36	0.0273	627.25
FBA-500ms	1.00	5.08	0.0341	827.62
Adaptive-100-250ms	1.00	4.71	0.0375	522.00
Adaptive-OrderFlow-100-250ms	1.00	5.93	0.0244	746.75
Adaptive-QueueLoad-100-250ms	1.00	4.88	0.0375	624.25
Policy-BurstAware-100-250ms	1.00	5.38	0.0355	637.62
FBA-250ms-Stress	1.00	5.35	0.0269	2057.00

Table 2: Aggregate market-quality and fairness-adjacent proxy metrics.

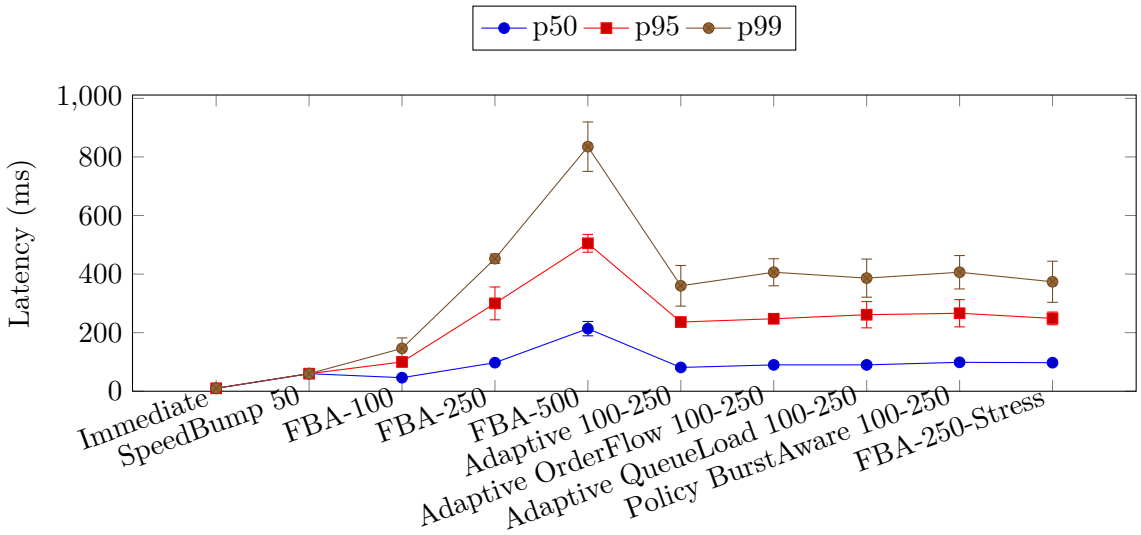

Figure 2: Latency profiles with 95% confidence intervals show the tail cost introduced by longer batch windows.

milliseconds and lowers queue advantage to 0.0244, but it gives back fills and arbitrage-profit performance. The queue-load adaptive variant settles at 209.29 milliseconds, improves price impact to 4.88, and keeps a lower p99 than the order-flow variant. The adapter-driven burst-aware controller saturates at a 250 millisecond mean window and lands between the adaptive heuristics and the fixed 250 millisecond batch on fills, queue advantage, and arbitrage-profit proxy. This does not prove fairness in a strong economic sense, but it does suggest that the benchmark is sensitive to mechanism-dependent priority effects.

The stress scenario increases throughput to 1769.25 orders per second and 900.50 fills per second. It also increases the arbitrage-profit proxy to 2057.00, which shows that higher activity does not automatically imply improved fairness-adjacent outcomes.

Across all $10 \times 8 = 80$ aggregate runs, the environment reports zero negative-balance violations and zero conservation breaches. This is the most important infrastructure-facing result at the current stage: benchmark comparisons are being made without giving up basic settlement safety.

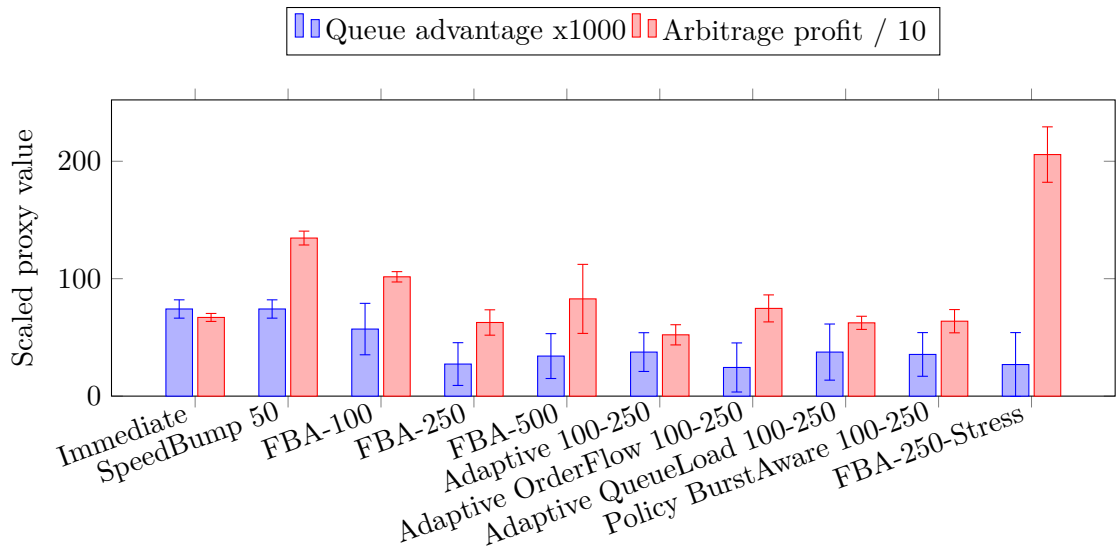

Figure 3: Fairness-adjacent proxy metrics with 95% confidence intervals show that the 250ms batch regime lowers queue advantage, while the stress case raises arbitrage-profit proxy.

8.3 Ablation Snapshot

We also generate a targeted ablation suite over seeds $[13, 17, 19, 23]$ for a stressed constrained 250 millisecond batch regime. Relaxing risk limits increases throughput from 1190.48 to 1770.24 orders per second and removes all risk rejections, but it also increases the arbitrage-profit proxy from 685.25 to 2198.50. Randomizing tie-breaks leaves throughput roughly unchanged while worsening p99 latency from 320.00 to 402.50 milliseconds. Disabling settlement checks does not change throughput in this setup, indicating that invariant enforcement is not the dominant bottleneck for the present benchmark scale.

9 Limitations

The current artifact is still short of a strong top-tier benchmark submission. The most important limitations are:

- only heuristic control, with no learned adaptive baseline
- proxy fairness metrics rather than richer behavioral or welfare metrics
- the adapter action space is still narrow and does not expose richer controls over pricing, order release timing, or agent behavior
- the current ablation suite is narrow and focuses on market-structure toggles rather than agent-model perturbations

10 Related Work

The benchmark is motivated directly by frequent-batch-auction market design [2]. Its benchmark-oriented framing is also aligned with the broader machine-learning practice of exposing reusable

environments rather than only single-task scripts [1]. The present contribution sits between those two traditions: it is neither a pure market-design theory paper nor a standard reinforcement-learning benchmark, but an infrastructure-aware environment that aims to make later learning or control work more credible.

11 Conclusion

This NeurIPS-track benchmark line establishes a reproducible, ledger-aware evaluation environment for market-design experiments without overwriting the original systems-paper line. The current results already show measurable latency and market-quality tradeoffs across immediate, speed-bump, adaptive, adapter-driven policy, and batch matching regimes, while preserving executable settlement invariants. The next stage is clear: add a learned adaptive controller, broaden the adapter into a richer policy interface, and introduce stronger behavioral and ablation studies so that the environment can support more serious benchmark or learning-oriented submissions.

References

[1] Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. Openai gym, 2016.

[2] Eric Budish, Peter Cramton, and John Shim. The high-frequency trading arms race: Frequent batch auctions as a market design response. *The Quarterly Journal of Economics*, 130(4):1547–1621, 2015.