A Ledger-Aware Benchmark for Market Infrastructure under Frequent Batch Auctions

Author Name
Institution
email@example.com

March 7, 2026

Abstract

We present a benchmark-oriented extension of a ledger-first market infrastructure prototype for studying market-design and execution behavior under settlement and risk constraints. The environment combines seedable agent-based order flow, configurable immediate, speed-bump, and frequent-batch-auction matching regimes, double-entry style account state transitions, and deterministic safety checks for conservation and non-negativity. We expose reproducible benchmark tasks that compare latency, throughput, spread, price impact, queue-priority advantage, and latency-arbitrage profit across market designs. In the current implementation, all generated scenarios preserve settlement invariants, while mechanism changes induce clear latency and market-quality tradeoffs. The result is a benchmark draft for market-infrastructure research that is separate from the original systems-paper line and is intended as a foundation for stronger benchmark-style evaluation.

1 Introduction

Electronic markets are shaped jointly by matching rules, latency structure, and settlement semantics. Market-design work on frequent batch auctions argues that short batching windows can reduce some latency-driven distortions in continuous markets [2]. In practice, many evaluation artifacts study matching rules in isolation and do not model the settlement layer, replay behavior, or safety checks that an actual market stack must satisfy.

This paper develops a benchmark-oriented layer on top of an existing ledger-first market-infrastructure prototype. It is not a replacement for the original systems paper. It is a separate evaluation track that asks a narrower question: can market-design tradeoffs be measured inside an environment that preserves explicit settlement invariants?

The current paper makes four concrete contributions:

1. It defines a ledger-aware benchmark environment in which market-design experiments are coupled to deterministic settlement checks.

2. It provides a seedable multi-agent order-flow generator with four agent classes, eleven benchmark scenarios, and explicit agent/workload ablation cases.

3. It reports single-seed and eight-seed aggregate benchmark results over throughput, latency, spread, price impact, queue-priority advantage, and arbitrage-profit proxies.

4. It exposes a step-wise `Reset/Step/Observe/Metrics` API, a minimal gym-style adapter with batch-window, risk-scale, and tie-break controls, two adapter-driven policy baselines, and ablation suites for both market-structure toggles and agent/workload perturbations.

2 Problem Setting

We study a discrete-time environment in which a market state is advanced one simulator step at a time. At step t, the state is represented as

$$s_t = (B_t, S_t, X_t, f_t),$$

where B_t and S_t denote the current buy and sell books, X_t is the vector of account states, and f_t is the latent fundamental used by the synthetic order-flow generator. A scenario additionally fixes a matching mode, a batch-window size, risk thresholds, and a population of agents.

The benchmark focuses on the following research questions:

1. How do immediate, speed-bump, adaptive, and frequent-batch-auction execution differ in latency and fill behavior?

2. Can fairness-adjacent signals such as queue-priority advantage and latency-arbitrage profit be measured reproducibly under deterministic settlement semantics?

3. Do settlement invariants hold under heterogeneous agents and stress scenarios?

3 Benchmark Environment

3.1 Core Components

The benchmark environment lives in the repository's `simulator/` package. It consists of scenario, order, fill, and result schemas, heterogeneous agent order generation, immediate, speed-bump, adaptive, and batch-clearing matching paths, deterministic settlement application, and metric accumulation.

The current environment also exposes a step-wise control surface with `Reset()`, `Observe()`, `Step()`, and `Metrics()`. On top of that surface, the repository includes a minimal adapter layer that returns reward-bearing timesteps and advertises per-scenario action support. In the current implementation, the adapter exposes three explicit control channels: batch-window override for adaptive scenarios, risk-limit scaling, and tie-break randomization for batch-style matching.

3.2 Determinism and Invariants

Each scenario is seeded. Given the same scenario configuration and seed, the environment is expected to produce the same accepted order counts, fill counts, and proxy metrics.

The benchmark currently enforces three operational invariants:

1. **Conservation**: total cash and total inventory across accounts remain constant over time.

2. **Non-negativity**: no account is allowed to end a fill application with negative cash or negative units.

3. **Deterministic replay**: seeded runs produce stable outputs for the same configuration.

4 Agent Models

The present environment includes four agent classes:

- **Market makers**: quote on both sides around the current fundamental with configurable width and intensity.
- **Latency arbitrageurs**: react to stale price information and attempt to exploit short-horizon dislocations.
- **Retail traders**: submit noisier buy or sell orders with wider dispersion around the current reference level.
- **Informed traders**: trade in the direction of the next-step fundamental when the generator exposes a directional signal.

These agents are intentionally simple. Their purpose is to generate heterogeneous but controlled activity so that matching rules can be compared under a repeatable workload. The stress scenario extends the default population with additional retail and arbitrage activity to increase order intensity.

5 Tasks and Baselines

The current benchmark suite contains eleven scenarios:

1. `Immediate-Surrogate`
2. `SpeedBump-50ms`
3. `FBA-100ms`
4. `FBA-250ms`
5. `FBA-500ms`
6. `Adaptive-100-250ms`
7. `Adaptive-OrderFlow-100-250ms`
8. `Adaptive-QueueLoad-100-250ms`
9. `Policy-BurstAware-100-250ms`
10. `Policy-LearnedLinear-100-250ms`
11. `FBA-250ms-Stress`

The benchmark now includes three heuristic adaptive-window baselines and two adapter-driven policy controllers. The burst-aware controller is hand-written. The learned linear controller is trained by seed-split model selection over a small linear threshold family and then executed through the same adapter action space.

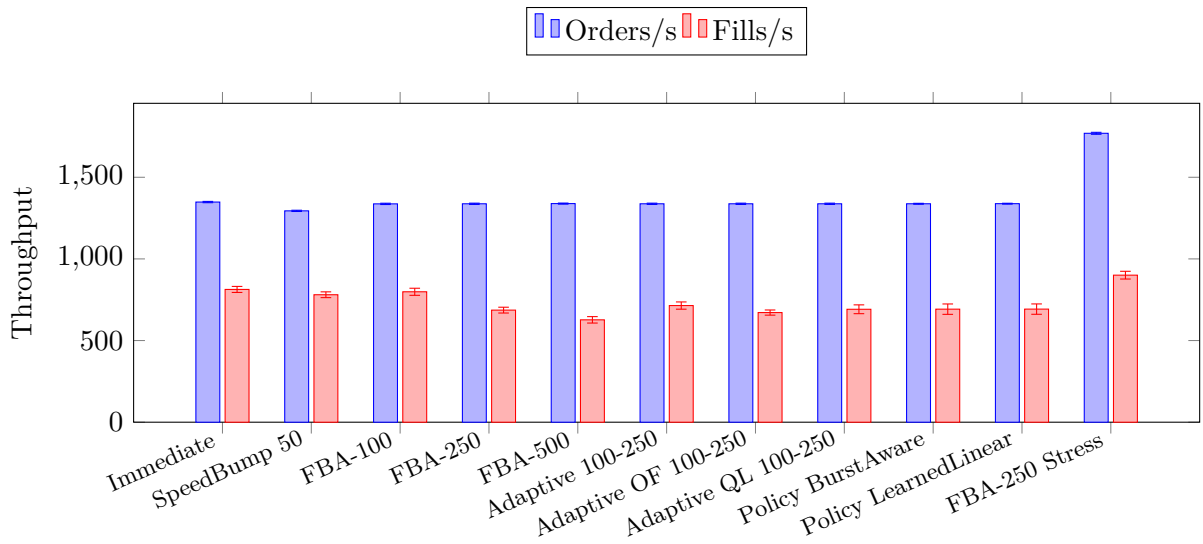

Figure 1: Aggregate throughput comparison with 95% confidence-interval error bars.

6 Metrics

The environment reports four groups of signals:

- **Systems metrics**: orders per second, fills per second, and p50/p95/p99 latency.
- **Market-quality metrics**: average spread and average price impact.
- **Fairness-adjacent proxies**: queue-priority advantage, latency-arbitrage profit, and execution dispersion.
- **Safety metrics**: negative-balance violations, conservation breaches, and risk rejections.

7 Experimental Setup

We report two result layers. The first is a single-seed reproducibility snapshot. The second is an eight-seed aggregate profile that is more appropriate for paper-facing discussion. The aggregate profile uses seeds

$$[7, 11, 19, 23, 29, 31, 37, 41].$$

All reported scenarios use the same simulator implementation and differ only in matching mode, batch-window size, seed, and population intensity. The current setup uses a discrete-time step duration of 10 milliseconds.

8 Results

8.1 Aggregate Summary

8.2 Observed Tradeoffs

Three measured observations are currently the most informative.

Scenario	Orders/s	Fills/s	p50 (ms)	p99 (ms)
Immediate-Surrogate	1348.2	813.1	10.0	10.0
SpeedBump-50ms	1294.3	780.6	60.0	60.0
FBA-100ms	1337.1	798.7	46.3	146.3
FBA-250ms	1337.6	686.4	97.5	452.5
FBA-500ms	1338.8	626.9	213.8	835.0
Adaptive-100-250ms	1337.6	714.3	81.3	360.0
Adaptive-OrderFlow-100-250ms	1337.6	671.4	90.0	406.3
Adaptive-QueueLoad-100-250ms	1337.6	691.6	90.0	386.3
Policy-BurstAware-100-250ms	1337.7	692.3	98.8	406.3
Policy-LearnedLinear-100-250ms	1338.3	692.6	92.5	333.8
FBA-250ms-Stress	1769.3	900.5	97.5	373.8

Table 1: Eight-seed aggregate summary. All runs report zero conservation breaches and zero negative-balance violations.

Scenario	Spread	Impact	Queue Adv.	Arb Profit
Immediate-Surrogate	1.98	3.38	0.0742	669.62
SpeedBump-50ms	1.96	5.96	0.0742	1346.25
FBA-100ms	1.00	5.96	0.0571	1015.75
FBA-250ms	1.00	5.36	0.0273	627.25
FBA-500ms	1.00	5.08	0.0341	827.62
Adaptive-100-250ms	1.00	4.71	0.0375	522.00
Adaptive-OrderFlow-100-250ms	1.00	5.93	0.0244	746.75
Adaptive-QueueLoad-100-250ms	1.00	4.88	0.0375	624.25
Policy-BurstAware-100-250ms	1.00	5.38	0.0355	637.62
Policy-LearnedLinear-100-250ms	1.00	5.51	0.0470	672.62
FBA-250ms-Stress	1.00	5.35	0.0269	2057.00

Table 2: Aggregate market-quality and fairness-adjacent proxy metrics.

First, immediate execution retains the lowest latency profile, with mean p50, p95, and p99 all equal to 10 milliseconds under the present discrete-time environment.

Second, the 50 millisecond speed-bump baseline introduces a fixed 60 millisecond latency profile and lowers throughput to 1294.30 orders per second while preserving the same queue-priority-advantage proxy observed under immediate execution. This makes it a useful intermediate baseline: latency rises, but batching-style fairness proxies do not automatically improve.

Third, moving from immediate execution to fixed batch windows creates a clear latency tradeoff without materially changing mean order throughput in the non-stress scenarios. The 100 millisecond batch already raises mean p99 latency to 146.25 milliseconds, while the 500 millisecond batch pushes mean p99 latency to 835.00 milliseconds.

Fourth, the 250 millisecond batch regime lowers the queue-priority-advantage proxy to 0.0273, below immediate execution (0.0742), the speed-bump baseline (0.0742), and the 100 millisecond batch regime (0.0571). The balanced adaptive heuristic settles around a mean window of 207.14 milliseconds and lowers the arbitrage-profit proxy to 522.00, below both fixed 100 and fixed 250 millisecond batching. The order-flow adaptive variant pushes to a higher mean window of 216.67

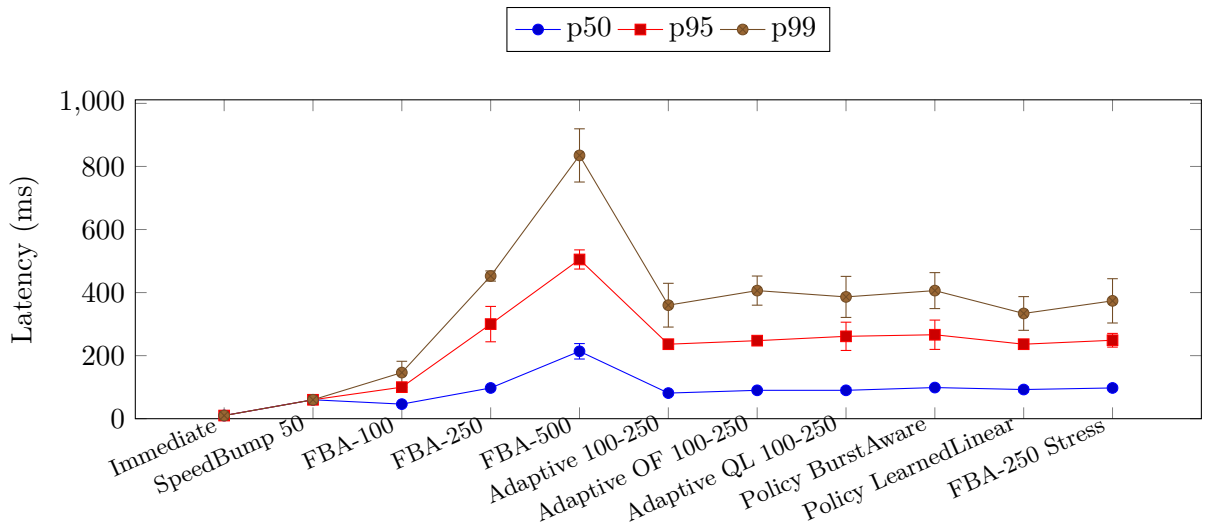

Figure 2: Latency profiles with 95% confidence intervals show the tail cost introduced by longer batch windows.

milliseconds and lowers queue advantage to 0.0244, but it gives back fills and arbitrage-profit performance. The queue-load adaptive variant settles at 209.29 milliseconds, improves price impact to 4.88, and keeps a lower p99 than the order-flow variant. The burst-aware controller saturates at a 250 millisecond mean window and lands between the adaptive heuristics and the fixed 250 millisecond batch on fills, queue advantage, and arbitrage-profit proxy. The learned linear controller shortens the mean window to 225.18 milliseconds and improves p99 to 333.75 milliseconds, but it does not improve the fairness-adjacent proxy scores relative to burst-aware. This is a useful benchmark result: the controller comparison now exposes a latency-versus-fairness tradeoff inside the adapter action space instead of only across built-in matching rules.

The stress scenario increases throughput to 1769.25 orders per second and 900.50 fills per second. It also increases the arbitrage-profit proxy to 2057.00, which shows that higher activity does not automatically imply improved fairness-adjacent outcomes.

Across all $11 \times 8 = 88$ aggregate runs, the environment reports zero negative-balance violations and zero conservation breaches.

8.3 Ablation Snapshots

We generate a targeted market-structure ablation suite over seeds $[13, 17, 19, 23]$ for a stressed constrained 250 millisecond batch regime. Relaxing risk limits increases throughput from 1190.48 to 1770.24 orders per second and removes all risk rejections, but it also increases the arbitrage-profit proxy from 685.25 to 2198.50. Randomizing tie-breaks leaves throughput roughly unchanged while worsening p99 latency from 320.00 to 402.50 milliseconds. Disabling settlement checks does not change throughput in this setup, indicating that invariant enforcement is not the dominant bottleneck for the present benchmark scale.

We also generate an agent/workload ablation suite over seeds $[43, 47, 53, 59]$. Removing arbitrageurs collapses arbitrage-profit proxy to zero and flips queue advantage negative. Removing informed traders reduces p99 latency to 335.00 milliseconds with little change in arbitrage-profit proxy. Increasing retail burst intensity pushes throughput above 2500 orders per second and raises p99 to 467.50 milliseconds. These results show that the benchmark is sensitive to population

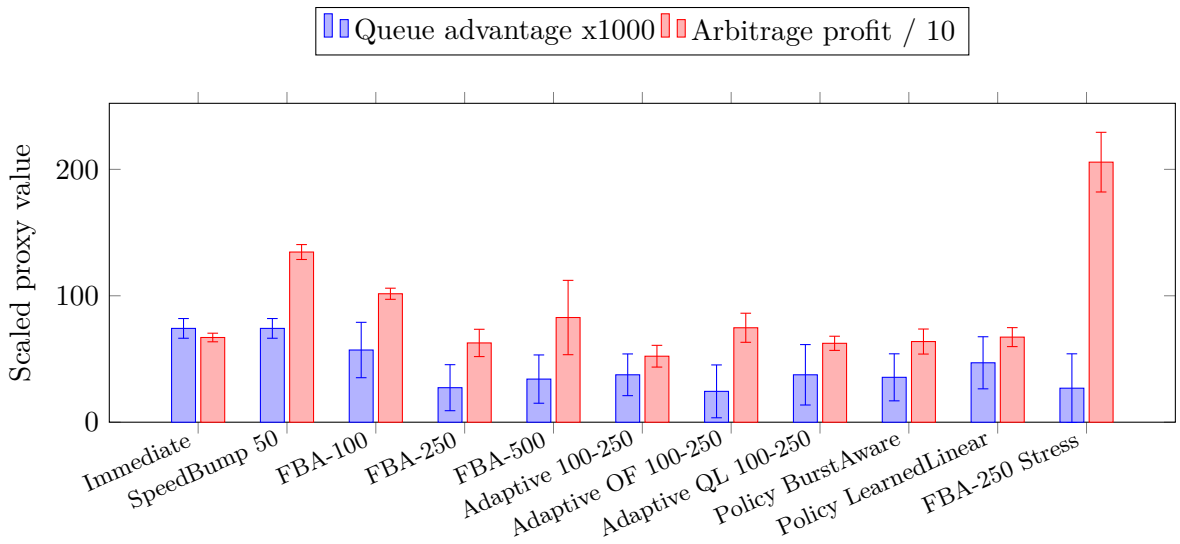

Figure 3: Fairness-adjacent proxy metrics with 95% confidence intervals. The learned controller improves latency tails but not fairness-adjacent proxy scores relative to the burst-aware policy.

composition and workload intensity, not only to the mechanism choice.

9 Limitations

The current artifact is still short of a strong top-tier benchmark submission. The most important limitations are:

- proxy fairness metrics rather than richer behavioral or welfare metrics
- the adapter action space is still narrow and does not expose richer controls over pricing, order release timing, or agent behavior
- the learned controller is still a tiny threshold family, not a serious policy-learning result
- the current agent/workload ablation suite is still small and does not yet cover richer market-maker or informed-flow parameter sweeps

10 Related Work

The benchmark is motivated directly by frequent-batch-auction market design [2]. Its benchmark-oriented framing is also aligned with the broader machine-learning practice of exposing reusable environments rather than only single-task scripts [1]. The present contribution sits between those two traditions: it is neither a pure market-design theory paper nor a standard reinforcement-learning benchmark, but an infrastructure-aware environment that aims to make later learning or control work more credible.

11 Conclusion

This NeurIPS-track benchmark line establishes a reproducible, ledger-aware evaluation environment for market-design experiments without overwriting the original systems-paper line. The current results already show measurable latency and market-quality tradeoffs across immediate, speed-bump, adaptive, adapter-driven policy, and batch matching regimes, while preserving executable settlement invariants. The next stage is clear: add a stronger learned controller, broaden the adapter into a richer policy interface, and introduce deeper behavioral and ablation studies so that the environment can support more serious benchmark or learning-oriented submissions.

References

[1] Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. Openai gym, 2016.

[2] Eric Budish, Peter Cramton, and John Shim. The high-frequency trading arms race: Frequent batch auctions as a market design response. *The Quarterly Journal of Economics*, 130(4):1547–1621, 2015.