A Ledger-Aware Benchmark for Market Infrastructure under Frequent Batch Auctions

Author Name
Institution
email@example.com

March 7, 2026

Abstract

We present a benchmark-oriented extension of a ledger-first market infrastructure prototype for studying market-design and execution behavior under settlement and risk constraints. The environment combines seedable agent-based order flow, configurable immediate, speed-bump, and frequent-batch-auction matching regimes, double-entry style account state transitions, and deterministic safety checks for conservation and non-negativity. We expose reproducible benchmark tasks that compare latency, throughput, spread, price impact, queue-priority advantage, and latency-arbitrage profit across market designs. In the current implementation, all generated scenarios preserve settlement invariants, while mechanism changes induce clear latency and market-quality tradeoffs. The result is a benchmark draft for market-infrastructure research that is separate from the original systems-paper line and is intended as a foundation for stronger benchmark-style evaluation.

1 Introduction

Electronic markets are shaped jointly by matching rules, latency structure, and settlement semantics. Market-design work on frequent batch auctions argues that short batching windows can reduce some latency-driven distortions in continuous markets [2]. In practice, many evaluation artifacts study matching rules in isolation and do not model the settlement layer, replay behavior, or safety checks that an actual market stack must satisfy.

This paper develops a benchmark-oriented layer on top of an existing ledger-first market-infrastructure prototype. It is not a replacement for the original systems paper. It is a separate evaluation track that asks a narrower question: can market-design tradeoffs be measured inside an environment that preserves explicit settlement invariants?

The current paper makes four concrete contributions:

1. It defines a ledger-aware benchmark environment in which market-design experiments are coupled to deterministic settlement checks.

2. It provides a seedable multi-agent order-flow generator with four agent classes, twelve benchmark scenarios, explicit agent/workload sweeps, a full parameter grid over arbitrage intensity and maker quote width, and a three-dimensional parameter cube over retail intensity, informed intensity, and maker quote width.

3. It reports single-seed and eight-seed aggregate benchmark results over throughput, latency, spread, price impact, queue-priority advantage, and arbitrage-profit proxies.

4. It exposes a step-wise `Reset/Step/Observe/Metrics` API, a gym-style adapter with five runtime control channels, three adapter-driven policy baselines, documented observation/action/metrics schemas, and ablation suites for both market-structure toggles and agent/workload perturbations.

2 Problem Setting

We study a discrete-time environment in which a market state is advanced one simulator step at a time. At step t, the state is represented as

$$s_t = (B_t, S_t, X_t, f_t),$$

where B_t and S_t denote the current buy and sell books, X_t is the vector of account states, and f_t is the latent fundamental used by the synthetic order-flow generator. A scenario additionally fixes a matching mode, a batch-window size, risk thresholds, and a population of agents.

The benchmark focuses on the following research questions:

1. How do immediate, speed-bump, adaptive, and frequent-batch-auction execution differ in latency and fill behavior?
2. Can fairness-adjacent signals such as queue-priority advantage and latency-arbitrage profit be measured reproducibly under deterministic settlement semantics?
3. Do settlement invariants hold under heterogeneous agents and stress scenarios?

3 Benchmark Environment

3.1 Core Components

The benchmark environment lives in the repository's `simulator/` package. It consists of scenario, order, fill, and result schemas, heterogeneous agent order generation, immediate, speed-bump, adaptive, and batch-clearing matching paths, deterministic settlement application, and metric accumulation.

The current environment also exposes a step-wise control surface with `Reset()`, `Observe()`, `Step()`, and `Metrics()`. On top of that surface, the repository includes an adapter layer that returns reward-bearing timesteps and advertises per-scenario action support. In the current implementation, the adapter exposes five explicit control channels: batch-window override for adaptive scenarios, risk-limit scaling, tie-break randomization, release-cadence control, and price-aggression bias.

3.2 Determinism and Invariants

Each scenario is seeded. Given the same scenario configuration and seed, the environment is expected to produce the same accepted order counts, fill counts, and proxy metrics.

The benchmark currently enforces three operational invariants:

1. **Conservation**: total cash and total inventory across accounts remain constant over time.
2. **Non-negativity**: no account is allowed to end a fill application with negative cash or negative units.
3. **Deterministic replay**: seeded runs produce stable outputs for the same configuration.

4 Agent Models

The present environment includes four agent classes:

- **Market makers**: quote on both sides around the current fundamental with configurable width and intensity.
- **Latency arbitrageurs**: react to stale price information and attempt to exploit short-horizon dislocations.
- **Retail traders**: submit noisier buy or sell orders with wider dispersion around the current reference level.
- **Informed traders**: trade in the direction of the next-step fundamental when the generator exposes a directional signal.

These agents are intentionally simple. Their purpose is to generate heterogeneous but controlled activity so that matching rules can be compared under a repeatable workload.

5 Tasks and Baselines

The current benchmark suite contains twelve scenarios:

1. `Immediate-Surrogate`
2. `SpeedBump-50ms`
3. `FBA-100ms`
4. `FBA-250ms`
5. `FBA-500ms`
6. `Adaptive-100-250ms`
7. `Adaptive-OrderFlow-100-250ms`
8. `Adaptive-QueueLoad-100-250ms`
9. `Policy-BurstAware-100-250ms`
10. `Policy-LearnedLinUCB-100-250ms`
11. `Policy-LearnedTinyMLP-100-250ms`
12. `FBA-250ms-Stress`

The benchmark now includes three heuristic adaptive-window baselines and three adapter-driven policy controllers. The burst-aware controller is hand-written. The learned baselines are a contextual linear bandit and a small two-layer policy network with burst-aware supervised warm-start and gradient-based policy updates over the same discrete action bundle set.

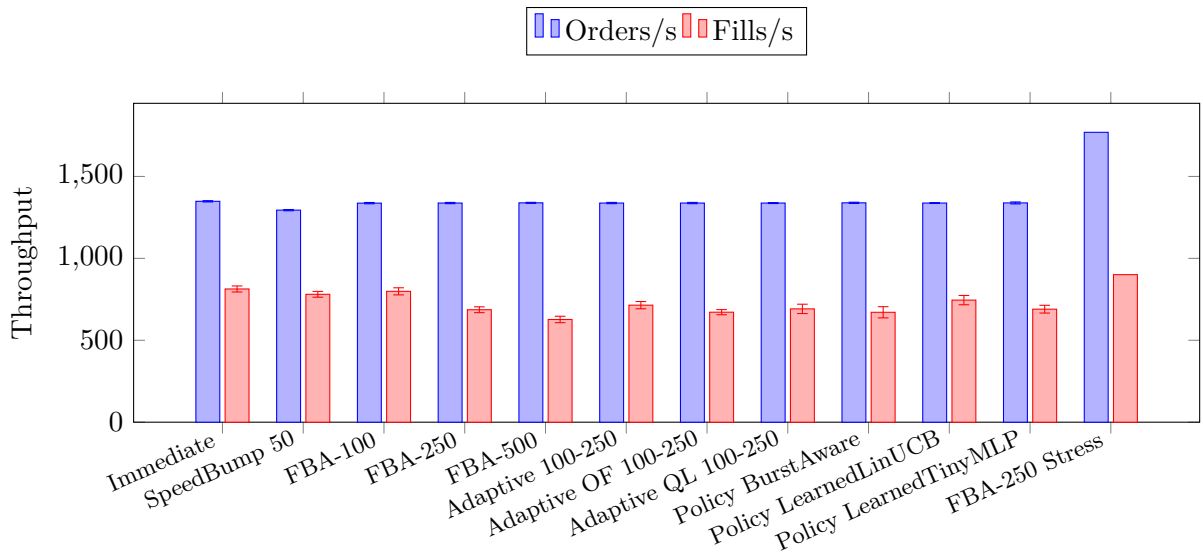

Figure 1: Aggregate throughput comparison with 95% confidence-interval error bars.

6 Metrics

The environment reports four groups of signals:

- **Systems metrics**: orders per second, fills per second, and p50/p95/p99 latency.
- **Market-quality metrics**: average spread and average price impact.
- **Fairness-adjacent proxies**: queue-priority advantage, latency-arbitrage profit, and execution dispersion.
- **Safety metrics**: negative-balance violations, conservation breaches, and risk rejections.

7 Experimental Setup

We report two result layers. The first is a single-seed reproducibility snapshot. The second is an eight-seed aggregate profile that is more appropriate for paper-facing discussion. The aggregate profile uses seeds

$$[7, 11, 19, 23, 29, 31, 37, 41].$$

All reported scenarios use the same simulator implementation and differ only in matching mode, batch-window size, seed, and population intensity. The current setup uses a discrete-time step duration of 10 milliseconds.

8 Results

8.1 Aggregate Summary

8.2 Observed Tradeoffs

Three measured observations are currently the most informative.

Scenario	Orders/s	Fills/s	p50 (ms)	p99 (ms)
Immediate-Surrogate	1348.2	813.1	10.0	10.0
SpeedBump-50ms	1294.3	780.6	60.0	60.0
FBA-100ms	1337.1	798.7	46.3	146.3
FBA-250ms	1337.6	686.4	97.5	452.5
FBA-500ms	1338.8	626.9	213.8	835.0
Adaptive-100-250ms	1337.6	714.3	81.3	360.0
Adaptive-OrderFlow-100-250ms	1337.6	671.4	90.0	406.3
Adaptive-QueueLoad-100-250ms	1337.6	691.6	90.0	386.3
Policy-BurstAware-100-250ms	1338.9	670.8	98.8	400.0
Policy-LearnedLinUCB-100-250ms	1337.6	745.2	47.5	158.8
Policy-LearnedTinyMLP-100-250ms	1338.4	689.8	58.8	305.0
FBA-250ms-Stress	1769.3	900.5	97.5	373.8

Table 1: Eight-seed aggregate summary. All runs report zero conservation breaches and zero negative-balance violations.

Scenario	Spread	Impact	Queue Adv.	Arb Profit
Immediate-Surrogate	1.98	3.38	0.0742	669.62
SpeedBump-50ms	1.96	5.96	0.0742	1346.25
FBA-100ms	1.00	5.96	0.0571	1015.75
FBA-250ms	1.00	5.36	0.0273	627.25
FBA-500ms	1.00	5.08	0.0341	827.62
Adaptive-100-250ms	1.00	4.71	0.0375	522.00
Adaptive-OrderFlow-100-250ms	1.00	5.93	0.0244	746.75
Adaptive-QueueLoad-100-250ms	1.00	4.88	0.0375	624.25
Policy-BurstAware-100-250ms	1.00	5.31	0.0305	621.00
Policy-LearnedLinUCB-100-250ms	1.00	5.94	0.0447	959.62
Policy-LearnedTinyMLP-100-250ms	1.00	8.86	0.0278	1229.00
FBA-250ms-Stress	1.00	5.35	0.0269	2057.00

Table 2: Aggregate market-quality and fairness-adjacent proxy metrics.

First, immediate execution retains the lowest latency profile, with mean p50, p95, and p99 all equal to 10 milliseconds under the present discrete-time environment.

Second, the 50 millisecond speed-bump baseline introduces a fixed 60 millisecond latency profile and lowers throughput to 1294.30 orders per second while preserving the same queue-priority-advantage proxy observed under immediate execution. This makes it a useful intermediate baseline: latency rises, but batching-style fairness proxies do not automatically improve.

Third, the 250 millisecond batch regime lowers the queue-priority-advantage proxy to 0.0273, below immediate execution (0.0742), the speed-bump baseline (0.0742), and the 100 millisecond batch regime (0.0571). The balanced adaptive heuristic settles around a mean window of 207.14 milliseconds and lowers the arbitrage-profit proxy to 522.00, below both fixed 100 and fixed 250 millisecond batching.

Fourth, the controller comparison now surfaces a sharper tradeoff. The burst-aware controller saturates at a 250 millisecond mean window and lands at lower queue-priority advantage (0.0305)

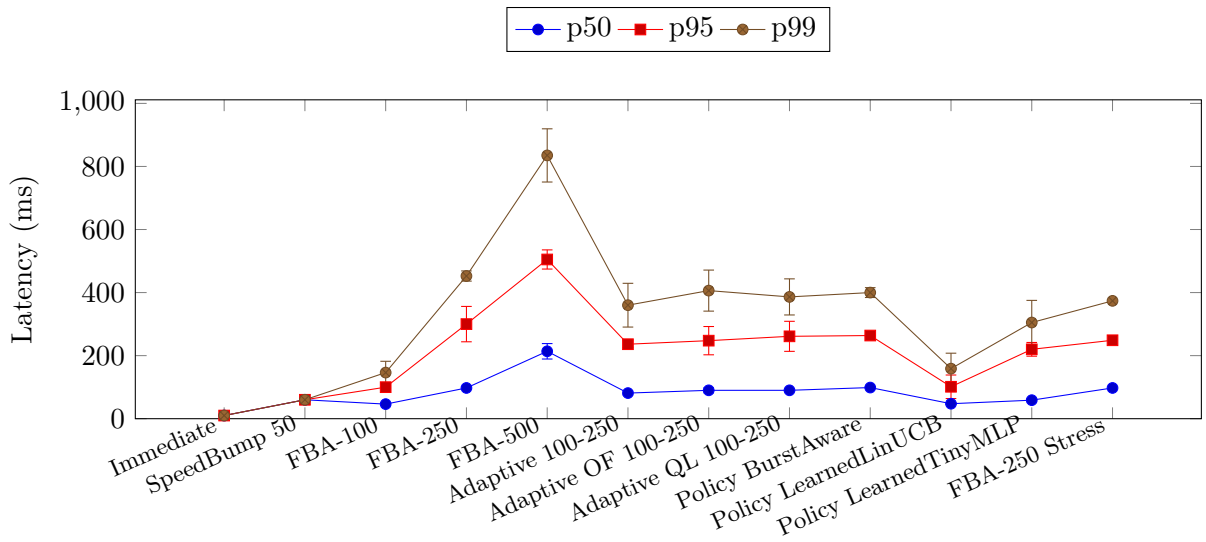

Figure 2: Latency profiles with 95% confidence intervals show the tail cost introduced by longer batch windows.

and lower arbitrage-profit proxy (621.00). The learned LinUCB controller collapses to a fast 100 millisecond control profile, improving fills to 745.24 per second and reducing p99 to 158.75 milliseconds, but giving back fairness-adjacent proxy scores (0.0447, 959.62). The gradient-trained TinyMLP controller shifts to a slower 200 millisecond mean window. It only modestly improves fills over burst-aware (689.78 versus 670.83) and cuts p99 to 305.00 milliseconds, but it also drives impact to 8.86 and arbitrage-profit proxy to 1229.00. The updated result is still useful for the benchmark line because it shows that stronger controller optimization does not automatically improve impact or fairness-adjacent outcomes.

The stress scenario increases throughput to 1769.25 orders per second and 900.50 fills per second. It also increases the arbitrage-profit proxy to 2057.00, which shows that higher activity does not automatically imply improved fairness-adjacent outcomes.

Across all $12 \times 8 = 96$ aggregate runs, the environment reports zero negative-balance violations and zero conservation breaches.

8.3 Ablation Snapshots

We generate a targeted market-structure ablation suite over seeds $[13, 17, 19, 23]$ for a stressed constrained 250 millisecond batch regime. Relaxing risk limits increases throughput from 1190.48 to 1770.24 orders per second and removes all risk rejections, but it also increases the arbitrage-profit proxy from 685.25 to 2198.50. Randomizing tie-breaks leaves throughput roughly unchanged while worsening p99 latency from 320.00 to 402.50 milliseconds. Disabling settlement checks does not change throughput in this setup, indicating that invariant enforcement is not the dominant bottleneck for the present benchmark scale.

We also generate an agent/workload sweep over seeds $[43, 47, 53, 59]$. Removing arbitrageurs collapses arbitrage-profit proxy to zero and flips queue advantage negative. Doubling arbitrage intensity pushes arbitrage-profit proxy to 1730.75. Widening market-maker quotes pushes p99 to 485.00 milliseconds and lowers fills to 595.63. Increasing informed intensity raises throughput modestly while leaving arbitrage-profit much closer to control than the arbitrage sweep. These results show that the benchmark is sensitive to population composition and workload intensity, not

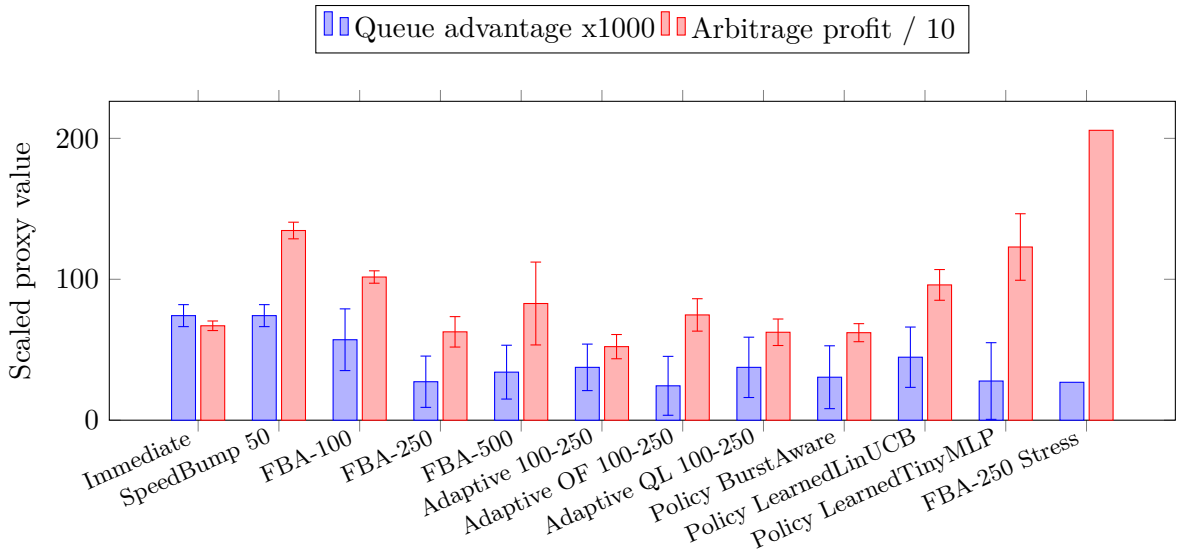

Figure 3: Fairness-adjacent proxy metrics with 95% confidence intervals. The gradient-trained TinyMLP controller lands closer to the batch-style queue proxy than LinUCB, but at substantially higher impact and arbitrage-profit proxy.

only to the mechanism choice.

Finally, we expose a full 4×3 parameter grid over arbitrageur intensity multipliers $\{0, 1, 2, 3\}$ and maker quote-width multipliers $\{1, 2, 3\}$, evaluated over seeds $[61, 67, 71, 73]$. Removing arbitrageurs collapses the arbitrage-profit proxy to zero across the entire first row and flips queue advantage strongly negative. Increasing arbitrage intensity raises throughput and arbitrage-profit proxy, reaching 2151.00 in the $(3, 3)$ cell. Wider maker quotes increase p99 tails and suppress fills, especially once arbitrage pressure is elevated.

We additionally expose a full $3 \times 3 \times 3$ parameter cube over retail-intensity multipliers $\{1, 2, 3\}$, informed-intensity multipliers $\{1, 2, 3\}$, and maker quote-width multipliers $\{1, 2, 3\}$, evaluated over seeds $[79, 83, 89, 97]$. Retail intensity is the dominant throughput lever in this cube, lifting orders per second from 1339.68 in the $(1, 1, 1)$ cell to 2132.54 in $(3, 1, 1)$. Informed intensity adds fills on top of that retail pressure, reaching 2292.86 orders per second and 1159.92 fills per second in $(3, 3, 1)$. Wider maker quotes suppress fills throughout the cube; for example, under $(3, 3, \cdot)$, fills fall from 1159.92 at maker width 1 to 956.75 at maker width 3. The cube keeps arbitrage-profit proxy in a much tighter band than the arbitrage-specific grid, which helps separate throughput stress from latency-arbitrage stress.

9 Limitations

The current artifact is still short of a strong top-tier benchmark submission. The most important limitations are:

- proxy fairness metrics rather than richer behavioral or welfare metrics
- the adapter action space is still narrow compared with a full exchange-control environment
- the strongest learned controller is still a small gradient-trained policy network rather than a richer offline-RL or larger policy-learning setup

- the current benchmark separates arbitrage-intensity sweeps from the retail / informed / maker cube instead of exposing a unified higher-dimensional stress surface

10 Related Work

The benchmark is motivated directly by frequent-batch-auction market design [2]. Its benchmark-oriented framing is also aligned with the broader machine-learning practice of exposing reusable environments rather than only single-task scripts [1]. The present contribution sits between those two traditions: it is neither a pure market-design theory paper nor a standard reinforcement-learning benchmark, but an infrastructure-aware environment that aims to make later learning or control work more credible.

11 Conclusion

This NeurIPS-track benchmark line establishes a reproducible, ledger-aware evaluation environment for market-design experiments without overwriting the original systems-paper line. The current results already show measurable latency and market-quality tradeoffs across immediate, speed-bump, adaptive, adapter-driven policy, and batch matching regimes, while preserving executable settlement invariants. With the addition of a gradient-trained policy network, a full arbitrage-intensity / maker-width parameter grid, and a three-dimensional retail / informed / maker cube, the benchmark now supports stronger controller comparisons and richer appendix-level sensitivity analysis. The next stage is to replace the current TinyMLP baseline with a stronger learned controller, broaden the action surface further, and merge the separate stress surfaces into a unified higher-dimensional sweep.

A Appendix Tables

Scenario	Orders/s	p99 (ms)	Arb Profit	Risk Rejects
Ablation-Control	1190.48	320.00	685.25	2922
Ablation-RelaxedRisk	1770.24	302.50	2198.50	0
Ablation-RandomTieBreak	1190.48	402.50	738.50	2908
Ablation-NoSettlementChecks	1190.48	320.00	685.25	2922

Table 3: Mechanism ablation appendix table.

References

[1] Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. Openai gym, 2016.

[2] Eric Budish, Peter Cramton, and John Shim. The high-frequency trading arms race: Frequent batch auctions as a market design response. *The Quarterly Journal of Economics*, 130(4):1547–1621, 2015.

Scenario	Orders/s	Fills/s	p99 (ms)	Arb Profit
AgentAblation-Control	1343.25	659.52	367.50	754.00
AgentAblation-NoArbitrageurs	1276.59	622.62	367.50	0.00
AgentAblation-NoInformed	1257.14	622.62	335.00	750.25
AgentSweep-RetailBurst	2532.94	1427.18	467.50	821.75
AgentSweep-ArbIntensityHigh	1409.92	718.45	302.50	1730.75
AgentSweep-InformedIntensityHigh	1426.79	724.01	360.00	685.50
AgentSweep-MakersWide	1343.25	595.63	485.00	754.50

Table 4: Agent and workload sweep appendix table.

Arb Mult.	Maker Width	Orders/s	p99 (ms)	Arb Profit
0	1	1272.42	390.00	0.00
1	1	1338.29	345.00	676.00
2	3	1404.17	462.50	1534.25
3	3	1470.04	462.50	2151.00

Table 5: Selected cells from the arbitrage-intensity / maker-width parameter grid.

Retail	Informed	Maker Width	Orders/s	Fills/s	Arb Profit
1	1	1	1339.68	650.60	508.50
1	3	1	1502.18	742.66	445.75
3	1	1	2132.54	1049.40	549.25
3	3	3	2292.86	956.75	461.75

Table 6: Selected cells from the retail / informed / maker parameter cube. Full slice figures are provided in the repository appendix.