Beyond Matching-Only Evaluation: Ledger-Aware Benchmarking for Learning in Market Infrastructure

Anonymous Authors

Abstract

Existing market-learning evaluations usually emphasize matching outcomes, latency, fills, or PnL-style summaries, and can therefore mis-rank controllers that would be unacceptable once settlement admissibility and welfare-sensitive outcomes are included. We introduce a ledger-aware benchmark that combines seedable agent-based order flow, mechanism variation, explicit conservation/non-negativity/replay invariants, a shared learning API, and a compact retail-welfare decomposition. The main empirical result is an evaluation failure: controller rankings and operating points change materially once settlement-aware and welfare-aware evaluation is used. Under the full protocol, controllers lie on a stable latency–welfare frontier: methods that improve tail latency and fill throughput often widen surplus-transfer gap, while more balanced methods recover retail-sensitive outcomes only by paying some latency cost. Value-style learners remain more stable on that frontier than online learners, which drift toward latency-favoring operating points unless checkpoint selection is treated as part of the protocol. These patterns persist under held-out regimes, reward perturbations, stress sweeps, and a first-pass Binance spot calibration envelope, suggesting that learning in market infrastructure requires settlement-aware evaluation rather than matching-only benchmarks.

1 Introduction

Electronic markets are shaped jointly by matching rules, latency structure, and settlement semantics. Yet many market-learning evaluations still study matching in isolation and summarize controller quality with latency, fills, or PnL-style proxies alone. For machine learning, that is not just an omission of detail. It creates an evaluation problem: a controller that looks strong in a matching-only simulator can still be misleading or unacceptable once settlement admissibility and welfare-sensitive outcomes are part of the protocol.

This paper develops a benchmark-oriented layer on top of an existing ledger-first market-infrastructure prototype. The benchmark is designed around a narrower but more ML-facing question than the original systems line: when mechanism choice, learned control, settlement invariants, and retail-sensitive measurement are placed in the same loop, do the conclusions about which controllers are good remain the same? We answer that question with a seedable environment that combines explicit conservation/non-negativity/replay checks, immediate and batched execution regimes, a shared adapter API, and a compact welfare decomposition.

The main result is an evaluation-failure story rather than a new RL update rule. Under the full protocol, controllers lie on a stable latency–welfare frontier: aggressive latency optimization improves p99 and fill throughput but tends to widen surplus-transfer gap, while more balanced controllers recover retail-sensitive outcomes only by paying some latency cost. Under the calibrated protocol, value-style learners remain more stable on that frontier than online learners, which drift toward latency-favoring operating points unless checkpoint selection is treated as part of the protocol.

Environment family	Synthetic/replay control	Mechanism variation	Settlement invariants	Learning API	Welfare summary
FinRL-Meta-style replay suites [6]	replay-heavy	limited	no	yes	no
ABIDES-Gym-style simulators [1]	yes	limited	no	yes	limited
JAX-LOB-style LOB simulators [4]	mostly synthetic	LOB-focused	no	limited	no
This benchmark	yes, seedable	immediate/batch/adaptive	yes	yes	yes

Table 1: Positioning relative to existing market-environment families. The distinctive feature here is not another simulator, but a learning-facing benchmark with explicit settlement invariants, mechanism variation, and retail-outcome summaries in the same loop.

Counterfactual controls make the necessity claim sharper. In the new ranking-shift artifact, moving from the full calibrated protocol to `matching_only` changes all four policy ranks and yields Kendall $\tau = -0.67$, whereas `no_settlement` leaves the same ranking unchanged with $\tau = 1.0$. Removing welfare terms from the reward also changes three of four ranks ($\tau = 0.33$). The point is therefore not that settlement dominates every aggregate metric. The point is that admissibility, welfare-sensitive outcomes, and protocol choice determine which conclusions can be trusted.

Why this benchmark is needed. Existing market environments usually emphasize one of three things: historical-replay reinforcement-learning pipelines, multi-agent market simulation, or high-throughput order-book execution. This benchmark is aimed at the gap between them. Mechanism effects, controller behavior, settlement semantics, and welfare-sensitive outcomes are often evaluated in separate layers even though production trading infrastructure couples them tightly. Real market replay also entangles mechanism effects with historical participant behavior and venue-specific structure. A seedable synthetic generator allows controlled stress testing of mechanism-controller interactions across arbitrage, retail, informed, and maker regimes that are difficult to isolate cleanly in historical data. The benchmark is therefore not just another market RL environment. Its purpose is to test whether simplifying the evaluation protocol changes what appears to be a good controller.

This paper makes four concrete contributions:

1. It defines a ledger-aware benchmark in which mechanism experiments, learned controllers, deterministic settlement checks, and welfare-sensitive outcomes are evaluated under one shared protocol.

2. It shows that simplified evaluation protocols can materially change controller rankings: `matching_only` reverses the calibrated policy ordering, while the full protocol and `no_settlement` agree on the same ranking.

3. It reports a unified learning story across offline and online controllers, including held-out regime evaluation, learning-curve dynamics, reward sensitivity, checkpoint-selection effects, and a standardized calibrated leaderboard in which value-style learners are more stable than online policy updates.

4. It strengthens the measurement and realism story with a four-dimensional stress surface, welfare-robustness summaries, and a first closed-loop market-data calibration pass in which the same latency–welfare tension persists.

2 Evaluation Under Settlement Constraints

We study a discrete-time environment in which a market state is advanced one simulator step at a time. At step t, the state is represented as

$$s_t = (B_t, S_t, X_t, f_t),$$

where B_t and S_t denote the buy and sell books, X_t is the vector of account states, and f_t is the latent synthetic fundamental used by the order-flow generator. A scenario additionally fixes a matching mode, a batch-window size, risk thresholds, and a population of agents.

The benchmark focuses on four concrete questions:

1. Do controller rankings change once settlement-aware and welfare-aware evaluation is used instead of matching-only summaries?
2. How do immediate, speed-bump, adaptive, and frequent-batch-auction execution differ in latency, fill behavior, and retail-sensitive outcomes?
3. Do offline/value-style and online learners exhibit different stability once checkpoint selection and held-out generalization are part of the protocol?
4. Do settlement invariants and welfare-sensitive metrics remain reproducible and interpretable under heterogeneous agents, stress scenarios, and calibration?

3 Benchmark Environment

3.1 Core Components

The benchmark couples seeded heterogeneous order flow, immediate/speed-bump/adaptive/batch execution, deterministic settlement application, and a metric layer that records systems, market-quality, and welfare-sensitive outcomes. Controllers interact with that environment through a common step-wise surface, `Reset()`, `Observe()`, `Step()`, and `Metrics()`, so heuristics, offline value learners, and online policy updates can be evaluated under one protocol.

The shared control bundle exposes five runtime channels: batch-window override for adaptive scenarios, risk-limit scaling, tie-break randomization, release-cadence control, and price-aggression bias [2]. The goal is not to exhaust the control design space. It is to provide a small but nontrivial action surface on which evaluation protocol, admissibility, and welfare-sensitive outcomes can all matter at once.

3.2 Determinism and Invariants

Each scenario is seeded. Given the same scenario configuration and seed, the environment is expected to produce the same accepted order counts, fill counts, and benchmark metrics. The benchmark currently enforces three operational invariants:

1. **Conservation**: total cash and total inventory across accounts remain constant over time.
2. **Non-negativity**: no account is allowed to end a fill application with negative cash or negative units.
3. **Deterministic replay**: seeded runs produce stable outputs for the same configuration.

4 Agent Models

The present environment includes four agent classes:

- **Market makers**: quote on both sides around the current fundamental with configurable width and intensity.
- **Latency arbitrageurs**: react to stale price information and attempt to exploit short-horizon dislocations.
- **Retail traders**: submit noisier buy or sell orders with wider dispersion around the current reference level.
- **Informed traders**: trade in the direction of the next-step fundamental when the generator exposes a directional signal.

These agents are intentionally simple. Their purpose is to generate heterogeneous but reproducible activity so that matching rules and controller choices can be compared under a repeatable workload. Richer strategic agents are a natural extension, but they are not required for the benchmark objective in its current form. To test whether the central result survives richer, state-dependent behavior, the artifact layer now also includes a strategic population with inventory-aware market makers, trend-reactive retail flow, signal-scaled informed traders, and dislocation-sensitive arbitrageurs.

5 Tasks and Baselines

The current benchmark suite contains fifteen scenarios:

1. `Immediate-Surrogate`
2. `SpeedBump-50ms`
3. `FBA-100ms`
4. `FBA-250ms`
5. `FBA-500ms`
6. `Adaptive-100-250ms`
7. `Adaptive-OrderFlow-100-250ms`
8. `Adaptive-QueueLoad-100-250ms`
9. `Policy-BurstAware-100-250ms`
10. `Policy-LearnedLinUCB-100-250ms`
11. `Policy-LearnedTinyMLP-100-250ms`
12. `Policy-LearnedOfflineContextual-100-250ms`
13. `Policy-LearnedFittedQ-100-250ms`

14. `Policy-LearnedOnlineDQN-100-250ms`

15. `FBA-250ms-Stress`

The controller suite now contains a hand-written burst-aware controller, a contextual linear bandit, a gradient-trained TinyMLP policy, an offline contextual value controller, an offline fitted-Q controller, and an online DQN-style controller. The offline controllers are trained from logged trajectories generated by burst-aware, LinUCB, TinyMLP, offline contextual, and random behavior policies. The online controller is trained directly through adapter interaction over the same discrete action bundle. These controllers are included as evaluative baselines rather than as claims of new learning methods. The benchmark question is whether different learning styles move toward different latency-welfare tradeoffs once infrastructure constraints are part of the environment definition. Beyond the paper-facing baseline suite, the artifact layer also adds a prioritized Double-DQN style online controller. It is reported as a learning-curve extension rather than folded into the main multiseed table because the goal is to test whether a stronger online-RL recipe changes the core latency-welfare tradeoff, not to claim a new best policy.

6 Metrics

The environment reports five groups of signals:

- **Systems metrics**: orders per second, fills per second, and p50/p95/p99 latency.
- **Market-quality metrics**: average spread and average price impact.
- **Fairness-adjacent proxies**: queue-priority advantage, latency-arbitrage profit, and execution dispersion.
- **Primary welfare decomposition**: retail surplus per traded unit, retail adverse-selection rate, and surplus-transfer gap.
- **Secondary diagnostics**: arbitrageur surplus per traded unit and welfare dispersion.
- **Safety metrics**: negative-balance violations, conservation breaches, and risk rejections.

The welfare metrics should be interpreted as controlled outcome proxies inside the benchmark environment rather than direct measurements of real-market welfare. Their purpose is to expose systematic tradeoffs between latency optimization and surplus transfer under a fixed synthetic fundamental, not to claim a universal welfare theorem for live markets.

7 Experimental Protocol

We separate exploratory sweeps from a paper-facing primary protocol. The primary protocol is the calibrated adaptive benchmark with fixed train / validation / held-out splits, a shared observation schema, a shared six-action control bundle, and explicit checkpoint rules. The primary reported axes are p99 latency and surplus-transfer gap; fills per second and retail surplus are secondary interpretation metrics, and the artifact layer additionally exposes a standardized leaderboard score that combines these terms under one fixed formula. The benchmark question is therefore not which learner wins after uncontrolled tuning. It is whether method conclusions survive once protocol, admissibility, and welfare-sensitive evaluation are fixed.

Controller	Function class	Training data / interaction	Main hyperparameters	Checkpoint rule
`FittedQ`	linear per-action values	offline trajectories from burst-aware, LinUCB, TinyMLP, offline contextual, and random policies on 6 train seeds	$\gamma = 0.97$, 8 Bellman iterations, least-squares refit each round, target clamp $[-25, 25]$	snapshots 0, 8 on held-out regimes; iteration 8 reported
`OnlineDQN`	TinyMLP Q-network (hidden dim 10)	online interaction on 6 train seeds with replay memory	$\gamma = 0.97$, 160 episodes, $\epsilon : 0.25 \to 0.03$, replay cap 6000, updates after 48 samples, 4 Q-updates per step, lr 0.0035, $L_2 = 10^{-4}$, target refresh every 250 env steps	checkpoints every 20 episodes; episode 160 reported
`DoubleDQN`	TinyMLP Q-network (hidden dim 14)	online interaction on 6 train seeds with prioritized replay	$\gamma = 0.97$, 200 episodes, $\epsilon : 0.20 \to 0.02$, replay cap 8000, updates after 64 samples, 6 Double-DQN updates per step, lr 0.0030, $L_2 = 10^{-4}$, Polyak target mix $\tau = 0.08$	checkpoints every 20 episodes evaluated on held-out regimes; curve reported separately

Table 2: Minimal training specification for the paper-facing learned controllers and the stronger Double-DQN artifact. These settings are intentionally lightweight: the benchmark claim is not a new RL algorithm, but that learning behavior can be measured under infrastructure constraints.

We report two result layers. The first is a single-seed reproducibility snapshot. The second is an eight-seed aggregate profile that is more appropriate for paper-facing discussion. The aggregate profile uses seeds

$$[7, 11, 19, 23, 29, 31, 37, 41].$$

The learned fitted-Q controller uses an explicit offline train/eval split. Logged trajectories are collected on training seeds $[181, 191, 193, 197, 199, 211]$ from burst-aware, LinUCB, TinyMLP, offline contextual, and random behavior policies. The fitted-Q loop runs for eight iterations with discount 0.97 over the shared discrete action bundle. The online DQN-style controller uses replay-based Q-learning over the same action bundle on training seeds $[307, 311, 313, 317, 331, 337]$, with checkpoints every 20 episodes through 160 episodes. Held-out evaluation uses seeds $[223, 227, 229, 233]$ on four unseen stress regimes defined in the artifact layer. All reported scenarios use the same simulator implementation and differ only in matching mode, batch-window size, seed, and population intensity. The current setup uses a discrete-time step duration of 10 milliseconds.

7.1 Learning Specification

The learned controllers share one adapter interface. Observations contain seven scalars: a bias term, queue depth, buy/sell imbalance, spread, pending-order count, risk-rejection count, and normalized episode progress. The shared action bundle exposes six discrete control choices over batch-window target, risk-limit scale, tie-break mode, release cadence, and price-aggression bias.

The per-step reward shared by all learned controllers is

$$\begin{aligned} r_t = {} & 1.0 \cdot \Delta\text{fills} - 0.25 \cdot \Delta\text{spread} - 1.5 \cdot \Delta\text{impact} - 12.0 \cdot |\Delta\text{queue}| \\ & - 0.005 \cdot \Delta\text{arb} + 18.0 \cdot \Delta\text{retail surplus} - 10.0 \cdot \Delta\text{retail adverse} - 2.0 \cdot \Delta\text{welfare dispersion} \\ & - 3.0 \cdot \max(\Delta\text{surplus gap}, 0) - 0.5 \cdot \Delta\text{risk rejects} - 10.0 \cdot \Delta\text{conservation breaches}. \end{aligned}$$

The calibrated protocol layer extends that setup with a versioned train / validation / held-out split on the calibrated adaptive benchmark. The artifact `simulator_calibrated_policy_protocol` uses train seeds $[1103, 1109, 1117, 1123]$, validation seeds $[1129, 1151]$, and held-out seeds $[1153, 1163, 1171, 1181]$. It reports `burst_aware`, `fitted_q`, a lightweight clipped PPO baseline [7], and an offline IQL-style baseline [5] on calibrated validation and held-out regimes. `ppo_clip`, `iql`, and `fitted_q` share the same observation schema, discrete action bundle, train / validation / held-out split, and evaluation budget; they differ only in learning rule, update schedule, and checkpoint selection. The purpose is not to claim a stronger RL method. It is to make the evaluation protocol explicit once the real-data calibration envelope is introduced.

On the reference machine used for artifact generation, the unified hypercube run (108 cells, 4 seeds per cell, 125 steps per run) completes 54,000 simulator steps in 19.42 seconds, or about 2.78×10^3 steps per second; the artifact layer records the same measurement in `simulator_runtime_profile`. Combining published per-cell throughput with episode duration yields an estimated wall-clock rate of 5.11×10^4 order events per second and 2.38×10^4 fills per second, which is sufficient for the lightweight offline and online learning experiments reported here.

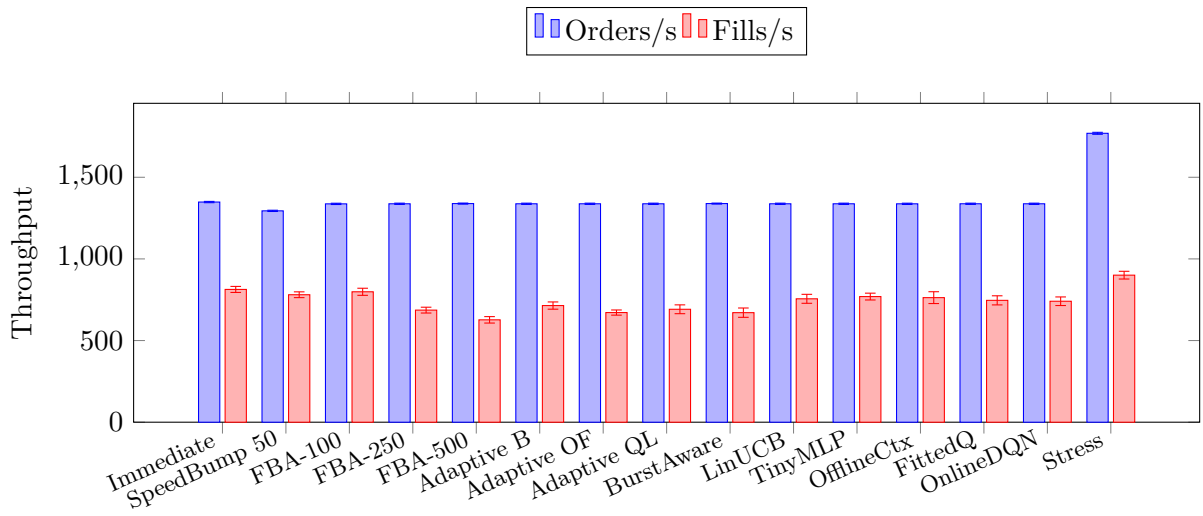

Figure 1: Aggregate throughput comparison with 95% confidence-interval error bars.

8 Results

The results are organized around four layers: the latency–welfare frontier itself, learning dynamics, protocol effects, and robustness under calibration and stress. The main claim is not only that a latency–welfare conflict exists. It is that simplified evaluation protocols can hide or distort that conflict, while the full protocol exposes stable differences between controller families.

Figure 2 is the most direct statement of benchmark necessity. The full protocol ranks `iql` first and `fitted_q` second, but the `matching_only` simplification reverses that ordering and promotes `ppo_clip` to the top position. The protocol question is therefore not cosmetic. Changing the evaluation loop changes what appears to be a good controller.

8.1 Aggregate Summary

8.2 Observed Tradeoffs

Four observations are currently the most informative.

First, immediate execution retains the lowest latency profile, but it also carries a high queue-advantage proxy and a large welfare gap. Under the current synthetic environment, low latency alone is not a good indicator of balanced market quality.

Second, the 50 millisecond speed-bump baseline raises latency without reducing queue advantage. It preserves the immediate-style queue signal (0.0742) while worsening the welfare gap to 4.1034. This makes it a useful negative control: latency rises, but batching-style fairness does not automatically follow.

Third, the fixed 250 millisecond batch regime lowers queue advantage to 0.0273 and compresses the welfare gap relative to immediate execution, but it gives up substantial tail latency. The balanced adaptive heuristic is the strongest fixed-logic mechanism on combined welfare terms: it keeps price impact low (4.71), reduces arbitrage-profit proxy to 522.00, and nearly closes the welfare gap (0.0278).

Fourth, the controller comparison now exposes a clearer tradeoff than the earlier draft. `LinUCB` remains strong on tail latency and fill throughput, but it does so by widening both queue advantage and the surplus-transfer gap. `FittedQ` and `OnlineDQN` both produce the best in-distribution p99

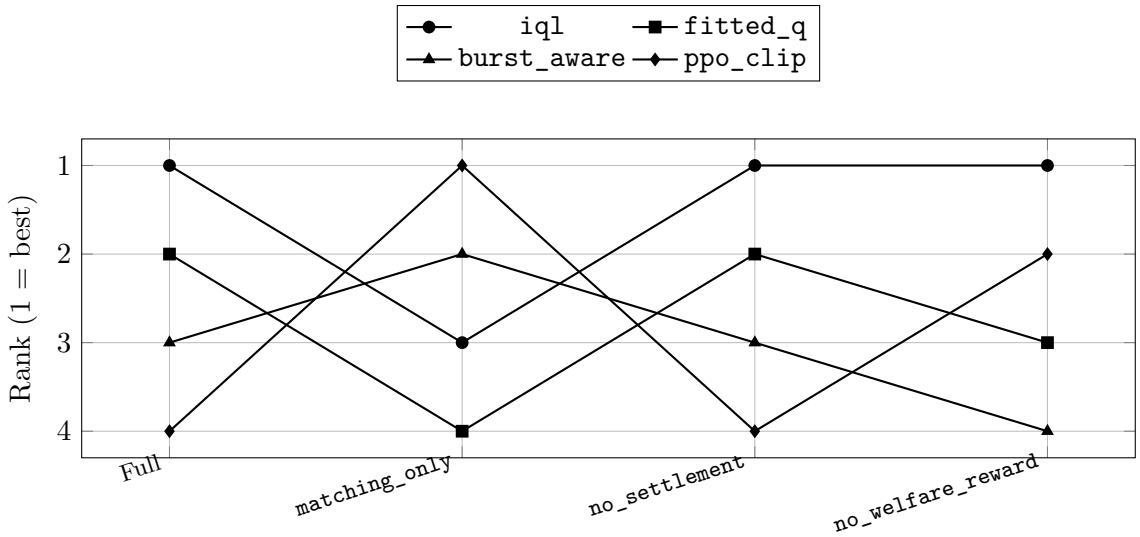

Figure 2: Method rankings change under evaluation protocol. The full ledger-aware protocol and `no_settlement` agree on the same ordering, while `matching_only` and `no_welfare_reward` materially reorder the policies.

tail (145.0 ms), giving explicit offline and online learning stories on top of the adapter API, but they still behave closer to `LinUCB` than to `OfflineContextual` on welfare transfer. `TinyMLP` improves fills further, but still leaves retail surplus negative and keeps arbitrage capture high. The offline contextual controller remains the strongest balanced learned baseline in the benchmark: compared with `LinUCB`, it gives up some tail latency but cuts price impact (4.94 versus 5.90), lowers queue advantage (0.0294 versus 0.0513), lowers arbitrage-profit proxy (771.25 versus 976.63), and shrinks the welfare gap (1.3769 versus 2.1694). This is already an evaluation effect: the methods that look strongest on p99 and fills are not the same methods that remain strongest once welfare-sensitive outcomes are treated as first-class metrics.

The stress scenario raises throughput to 1769.25 orders per second and 900.50 fills per second, but it also raises arbitrage capture to 2057.00. Higher activity does not automatically imply better fairness or welfare outcomes.

Across all $15 \times 8 = 120$ aggregate runs, the environment reports zero negative-balance violations and zero conservation breaches.

8.3 Held-Out Regime Generalization

Held-out results use unseen seeds $[223, 227, 229, 233]$ and four unseen regime combinations. The policy summary is shown in Table 5.

The held-out picture is materially different from the in-distribution aggregate. `FittedQ` lowers both p99 and welfare gap relative to `LinUCB` in the high-arbitrage wide-maker, retail-burst, and composite-stress held-out regimes. In the informed-wide regime, it gives up some tail latency (182.5 versus 167.5 ms) but still lowers welfare gap (2.2994 versus 2.3765). `OfflineContextual` remains the most welfare-balanced learned baseline overall, but it is materially slower than both `LinUCB` and `FittedQ` on held-out tails. `OnlineDQN` sits between them: it improves held-out fills to 986.61, keeps p99 at 172.50 ms, and lowers the held-out welfare gap to 2.2189, which is better than both `LinUCB` and `FittedQ` on the held-out aggregate.

Scenario	Orders/s	Fills/s	p50 (ms)	p99 (ms)
Immediate-Surrogate	1348.2	813.1	10.0	10.0
SpeedBump-50ms	1294.3	780.6	60.0	60.0
FBA-100ms	1337.1	798.7	46.3	146.3
FBA-250ms	1337.6	686.4	97.5	452.5
FBA-500ms	1338.8	626.9	213.8	835.0
Adaptive-100-250ms	1337.6	714.3	81.3	360.0
Policy-BurstAware	1338.9	670.8	98.8	400.0
Policy-LearnedLinUCB	1337.6	755.7	48.8	155.0
Policy-LearnedTinyMLP	1337.6	769.3	62.5	221.3
Policy-LearnedOfflineContextual	1337.4	762.8	65.0	215.0
Policy-LearnedFittedQ	1337.7	746.2	46.3	145.0
Policy-LearnedOnlineDQN	1337.6	740.8	46.3	145.0
FBA-250ms-Stress	1769.3	900.5	97.5	373.8

Table 3: Eight-seed aggregate summary. All runs report zero conservation breaches and zero negative-balance violations.

Scenario	Impact	Queue Adv.	Retail Surplus	Welfare Gap
Immediate-Surrogate	3.38	0.0742	-0.3710	2.0430
FBA-250ms	5.36	0.0273	-0.3493	0.8896
Adaptive-100-250ms	4.71	0.0375	-0.1508	0.0278
Policy-BurstAware	5.31	0.0305	-0.2258	0.7579
Policy-LearnedLinUCB	5.90	0.0513	0.0795	2.1694
Policy-LearnedTinyMLP	5.22	0.0336	-0.3128	1.9719
Policy-LearnedOfflineContextual	4.94	0.0294	-0.1090	1.3769
Policy-LearnedFittedQ	5.88	0.0451	0.0742	2.2036
Policy-LearnedOnlineDQN	5.92	0.0431	0.0662	2.2472
FBA-250ms-Stress	5.35	0.0269	-0.1494	0.3805

Table 4: Aggregate market-quality, welfare, and fairness-adjacent summary metrics.

8.4 Fitted-Q Learning Curve

The fitted-Q controller is no longer just a point baseline. We also log its held-out behavior over Bellman-update iterations. The untrained snapshot begins at a welfare gap of 3.2117 ± 0.5538 and p99 latency of 341.25 ± 22.97 ms. After the first Bellman update, the held-out welfare gap falls to 2.0790 ± 0.4767 and p99 falls to 198.75 ± 25.39 ms. By iteration 8, Bellman MSE declines from 45.1571 to 6.5753, held-out p99 reaches 155.62 ± 12.84 ms, and the welfare gap settles at 2.4226 ± 0.5400. The training story is therefore not monotone on every objective: later iterations improve Bellman error and tail latency, but they give back part of the earliest welfare improvement.

8.5 Online DQN Learning Curve

The online DQN-style controller adds a second, explicitly online training story. The untrained checkpoint begins at held-out p99 latency of 200.00 ± 25.46 ms and welfare gap of 1.5898 ± 0.3869. By episode 20, held-out p99 falls to 155.62 ± 12.84 ms, while the welfare gap rises to 2.4226 ± 0.5400.

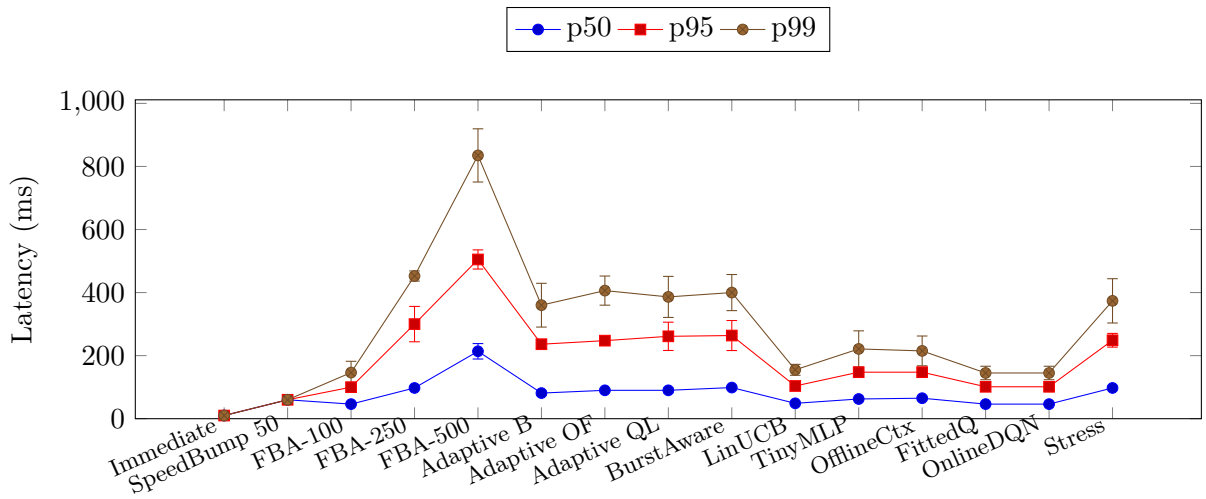

Figure 3: Latency profiles with 95% confidence intervals.

Policy	Fills/s	p99 (ms)	Retail Surplus	Welfare Gap
BurstAware	834.08	376.88	-0.4042	1.6286
LinUCB	978.52	171.25	0.1215	2.4466
OfflineContextual	961.61	275.62	-0.0732	1.9535
FittedQ	946.78	159.38	0.1110	2.3158
OnlineDQN	986.61	172.50	0.0937	2.2189

Table 5: Held-out policy summary over four unseen regimes.

Later checkpoints remain on the same plateau. The online controller therefore converges quickly toward a latency-favoring regime. Under the current reward structure, latency improvement is easier for the policy to exploit than retail-welfare preservation.

8.6 Prioritized Double-DQN Extension

The stronger online-RL artifact uses prioritized replay, Double-DQN targets, and Polyak target-network mixing. It does not change the benchmark question, but it does add a materially stronger online-learning recipe than the plain DQN baseline. The untrained checkpoint starts at 875.55 ± 145.75 fills per second, 336.25 ± 37.31 ms p99 latency, and welfare gap 3.2567 ± 0.6388. By episode 80, the controller reaches 1079.32 ± 183.72 fills per second with p99 latency 177.50 ± 21.88 ms and welfare gap 1.6813 ± 0.4221. By the final checkpoint at episode 200, p99 falls further to 161.88 ± 14.82 ms, but the welfare gap widens back to 2.5926 ± 0.5294. The final iterate therefore improves p99 by only 15.62 ms relative to episode 80, while giving back 0.9113 welfare-gap units and 145.64 fills per second. The important result is not domination over the simpler DQN baseline. It is that a stronger online learner exposes a sharper checkpoint-selection tradeoff: later optimization can still pull the policy back toward a latency-favoring regime, so checkpoint selection must be treated as part of the benchmark protocol rather than as an afterthought.

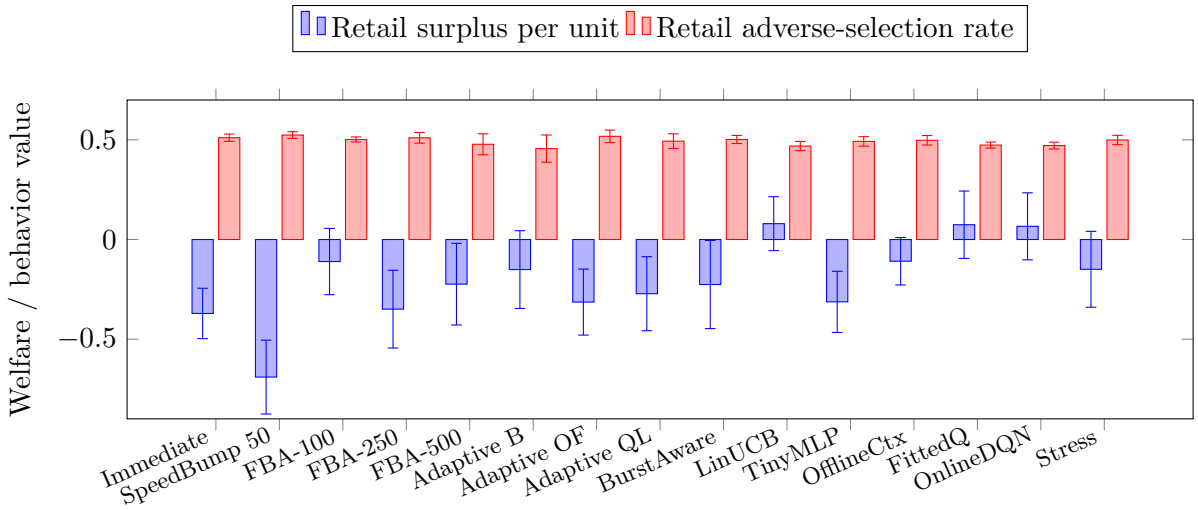

Figure 4: Welfare and behavior signals with 95% confidence intervals. Controllers that improve latency do not automatically improve retail outcomes.

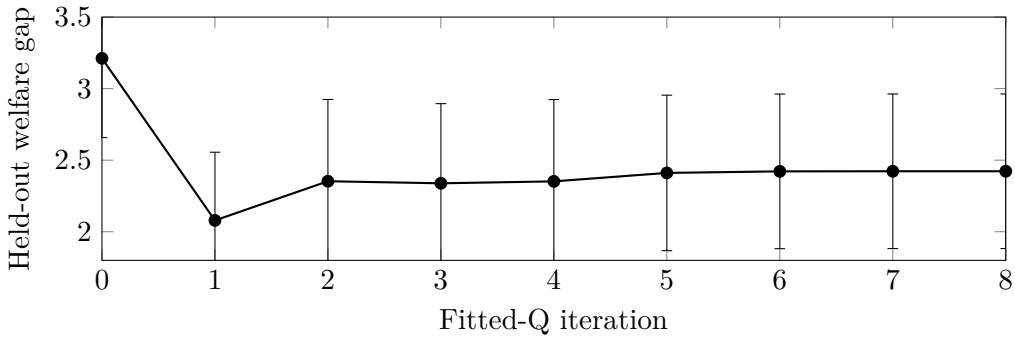

Figure 5: Held-out fitted-Q learning curve. The first Bellman update captures most of the welfare-gap gain; later iterations continue to reduce Bellman error and tail latency but do not preserve that full early welfare improvement.

8.7 Reward Sensitivity

To test whether the online-learning result depends on one hand-tuned reward vector, we retrain the online DQN controller under three deliberately separated reward profiles: the paper default, a latency-heavy profile, and a welfare-heavy profile. The final held-out operating point remains on the same plateau in all three cases. Default reward weights produce held-out p99 latency of 155.62 ± 12.84 ms and welfare gap 2.4226 ± 0.5400. The latency-heavy and welfare-heavy profiles produce the same held-out p99 and welfare-gap pair within the reported confidence intervals. The training reward scale changes materially, but the held-out controller behavior does not. For this benchmark, that is a useful robustness result: the latency-welfare tension is not only a consequence of one reward-weight vector. Under the current action bundle and held-out regimes, the online policy still converges toward the same latency-favoring operating region.

Profile	Mean train reward	Held-out p99 (ms)	Held-out welfare gap
Default	923.96	155.62	2.4226
Latency-heavy	1899.13	155.62	2.4226
Welfare-heavy	-108.57	155.62	2.4226

Table 6: Online DQN reward-sensitivity summary on held-out regimes. Large reward-weight perturbations change train reward scale but not the final held-out operating point.

8.8 First Closed-Loop Calibration Pass

The benchmark now includes a first closed-loop calibration pass against real public market data rather than only a synthetic generator. A stylized-facts envelope from an 8-symbol Binance spot slice is used to compare a baseline synthetic scenario against a tuned synthetic scenario.

The first-pass result is directionally strong even though it is not yet a final realism claim. The tuned generator moves mean spread from 326.41 bps to 12.83 bps, near the empirical upper envelope of 11.22 bps, and moves the top impact bucket from 264.41 bps to 8.47 bps. It falls inside the empirical range on 7/13 tracked summary metrics. The remaining misses are concentrated in order-sign autocorrelation, inter-arrival timing, one return-clustering term, and a subset of ask-side depth-shape statistics.

More importantly for the benchmark claim, the calibrated benchmark rerun in `simulator_calibrated_benchma` keeps all calibrated scenarios at zero-breach while preserving the same qualitative latency–welfare tension as the base benchmark:

- `Calibrated-Immediate-Surrogate`: 34.84 ± 0.01 orders/s, welfare gap 4.3321 ± 0.1398
- `Calibrated-FBA-2s`: 34.77 ± 0.01 orders/s, welfare gap 3.3927 ± 0.1593
- `Calibrated-Policy-LearnedFittedQ-1-3s`: 34.77 ± 0.01 orders/s, welfare gap 3.1301 ± 0.1299

The realism gap is therefore narrower than in earlier versions of the benchmark, even though this remains a first-pass calibration rather than a high-fidelity market simulator. The central latency–welfare tension does not disappear once the synthetic generator is retuned toward a market-data envelope.

8.9 Protocol Effects and Counterfactual Controls

The raw market-data bundle is versioned with manifest hashes and per-file checksums in the artifact package. This does not turn simulator outputs into real-market results, but it does make the calibration input chain auditable and reproducible.

The calibrated-learning protocol now reports explicit train / validation / held-out splits rather than only a single tuned rerun. On held-out calibrated regimes, `burst_aware` reaches 47.16 ± 3.43 fills per second with welfare gap 5.8322 ± 0.1675, `fitted_q` reaches 56.86 ± 1.89 fills per second with welfare gap 3.9602 ± 0.2672, `ppo_clip` reaches 40.54 ± 1.85 fills per second with welfare gap 6.6590 ± 0.3586, and `iql` reaches 58.86 ± 1.72 fills per second with welfare gap 3.9602 ± 0.2590. The strongest calibrated held-out tradeoff therefore comes from the IQL-style baseline rather than the simple PPO-style controller. A paired sign-flip comparison on the calibrated held-out regimes shows that `iql` improves both welfare gap and p99 latency relative to `ppo_clip` with exact $p = 3.1 \times 10^{-5}$ under the current protocol. The standardized calibrated leaderboard in `simulator_policy_leaderboard` therefore ranks `iql` first and `fitted_q` second under one fixed score combining fills, p99, retail

Protocol variant	Kendall τ vs control	Rank changes	Frontier overlap
`matching_only`	-0.67	4 / 4	1
`no_settlement`	1.00	0 / 4	2
`no_welfare_reward`	0.33	3 / 4	1

Table 7: Controller ordering changes under simplified evaluation protocols. The strongest protocol effect comes from reducing the environment to `matching_only`; removing settlement alone does not change the ordering in the current aggregate setup.

Frontier point	p99 (ms)	Welfare gap	Fills/s
Immediate-Surrogate	10.00	2.0430	813.13
Policy-LearnedOfflineContextual	215.00	1.3769	762.80
Adaptive-100-250ms	360.00	0.0278	714.29

Table 8: Undominated points on the multiseed controller Pareto surface.

surplus, adverse selection, and surplus-transfer gap. The stronger Double-DQN artifact then adds a second protocol finding: once online learning is strong enough, checkpoint selection itself becomes part of the benchmark because intermediate policies can dominate final ones on the latency-welfare frontier. Under the calibrated envelope, value-style learners are therefore more stable on the latency-welfare frontier, whereas online learners are more prone to drift toward latency-favoring operating points unless checkpoint selection is treated as part of the protocol.

The counterfactual-control artifact adds three negative controls around the calibrated benchmark. In `matching_only`, p99 latency collapses to 1000 ms but retail surplus stays negative and welfare gap remains large (4.5772 to 5.2132 across published policies). More importantly, the policy ranking changes materially: relative to the full control protocol, all four policy ranks move and Kendall $\tau = -0.67$. In `no_settlement`, the mechanism and welfare measurements remain unchanged while settlement application is removed, and the same policy ranking is preserved with Kendall $\tau = 1.0$. In `no_welfare_reward`, the PPO-style controller improves to 54.49 ± 4.34 fills per second and 3312.50 ± 227.12 ms p99, but still leaves a substantial welfare gap at 3.3874 ± 0.1272; the protocol ranking also changes, with three of four policy ranks moving and Kendall $\tau = 0.33$. These controls sharpen the benchmark claim: the main tension does not disappear when reward weights are altered, and controller conclusions can change qualitatively when the infrastructure setting is reduced to matching-only execution. The `no_settlement` control should not be read as evidence that settlement is unimportant. It should be read more narrowly: in the present aggregate setup, mechanism-level summary statistics are dominated by the matching/controller layer, while settlement still defines correctness, admissibility, and the trust boundary under which a learned controller would be acceptable in infrastructure use.

8.10 Controller Pareto Frontier

The multiseed controller frontier makes the same tradeoff visible on one surface. When minimizing p99 latency and surplus-transfer gap jointly, three points remain undominated: immediate execution, the balanced adaptive heuristic, and the offline contextual controller. Immediate execution is fastest but welfare-weak; the adaptive heuristic nearly closes the welfare gap but gives up substantial tail latency; offline contextual is the strongest balanced learned point.

Factor	Δ Orders/s	Δ Retail Surplus	Δ Retail Adverse	Δ Welfare Gap
Arbitrage $0 \to 3$	176.26	-0.1804	0.0117	1.2099
Retail $1 \to 3$	780.94	0.0772	0.0018	0.0781
Informed $1 \to 3$	150.91	-0.1986	0.0131	0.0876
Maker width $1 \to 3$	0.00	-0.1250	0.0085	0.2653

Table 9: Compact high-low contrasts from the unified hypercube.

8.11 Appendix-Supported Robustness

Secondary robustness details now live in the appendix rather than in the main narrative. The appendix retains mechanism ablations, agent/workload sweeps, and strategic-agent snapshots. Their role is supportive rather than central: they show that the main claim is not tied to one frozen population mix, one risk setting, or one single market-maker specification. Across those appendix-supported checks, the same directional pattern persists: relaxing risk raises throughput but also arbitrage capture, retail burst loads activity without neutralizing welfare transfer, and richer strategic agents preserve the same arbitrage-driven deterioration in retail outcome. The new welfare-robustness artifact makes the measurement story more explicit: across base and strategic populations, surplus-transfer gap retains moderate positive correlation with queue advantage and price impact, while the welfare-gap ranking remains substantially stable across suites (Kendall $\tau = 0.6$).

Finally, we expose a unified $4 \times 3 \times 3 \times 3$ hypercube over arbitrageur intensity $\{0, 1, 2, 3\}$, retail intensity $\{1, 2, 3\}$, informed intensity $\{1, 2, 3\}$, and maker quote width $\{1, 2, 3\}$. Two raw cells still illustrate the environment directly. At $(\text{arb} = 0, \text{retail} = 1, \text{informed} = 2, \text{maker} = 1)$, the environment reports 1350.60 orders per second, zero arbitrage-profit proxy, and a welfare gap of 0.2376. At $(\text{arb} = 3, \text{retail} = 1, \text{informed} = 2, \text{maker} = 1)$, throughput rises to 1542.86 and p99 falls to 362.50 milliseconds, but arbitrage-profit proxy jumps to 1982.00 and the welfare gap widens to 1.2839.

This compact summary removes the need to interpret the hypercube only through slice families. Rather than reporting slice plots alone, the benchmark summarizes the stress surface with response-surface fits and factor contrasts that expose the dominant drivers of welfare transfer. Arbitrage intensity consistently emerges as the dominant driver of welfare-gap expansion, retail intensity is mostly a throughput lever, and wider maker quotes worsen retail outcome without moving aggregate activity in this setup. Averaged over informed intensity and maker width, moving from arb = 0 to arb = 3 widens the welfare gap by 1.1748, 1.2262, and 1.2287 at retail intensities $x1$, $x2$, and $x3$, respectively. This effect persists across retail intensities. Higher retail flow does not neutralize transfer-to-arbitrageur.

The fitted response surface compresses the same hypercube into standardized main effects and pairwise interactions. For surplus-transfer gap, the fit reaches $R^2 = 0.3495$, with arbitrage intensity dominating partial variance explained (0.3157) and maker quote width the next largest effect (0.0291). For retail surplus per unit, the fit reaches $R^2 = 0.7061$, where informed intensity (0.3121) and arbitrage intensity (0.2374) dominate the effect ranking. This moves the benchmark beyond purely descriptive slice analysis.

8.12 Robustness to Agent Assumptions

The benchmark does not claim that its agents are complete models of strategic market behavior. The narrower question is whether the main tradeoffs survive changes in population mix and

Metric	R^2	Top partial effects
Surplus-transfer gap	0.3495	arb 0.3157, maker 0.0291
Retail surplus per unit	0.7061	informed 0.3121, arb 0.2374

Table 10: Response-surface summary over the unified hypercube.

workload composition. In the current artifact, they do. Removing arbitrageurs collapses arbitrage-profit proxy to zero and flips queue advantage negative. Raising arbitrage intensity in the agent sweep lifts arbitrage-profit proxy to 1730.75, while the hypercube summary shows welfare-gap expansion of about +1.2 when arbitrage intensity moves from 0 to 3. Retail burst nearly doubles throughput to 2532.94 orders per second but does not neutralize transfer-to-arbitrageur in the high-arbitrage slices. Wider maker quotes worsen retail outcome and p99 latency without meaningfully increasing aggregate activity. The richer strategic population preserves the same directional story: `Strategic-HighArb` worsens retail outcome relative to `Strategic-Control`, while `Strategic-RetailBurst` mainly raises throughput and fills rather than closing the welfare gap. These checks do not make the agents realistic in a behavioral-economics sense, but they do show that the central finding is not an artifact of one frozen population mix.

9 Scope and Limitations

The benchmark has clear remaining limitations. The most important ones are:

- the learned-controller family is still lightweight and discrete-action; this is a deliberate scope choice because the paper studies learning under infrastructure constraints rather than algorithmic novelty
- welfare metrics are derived from a synthetic fundamental rather than real market replay, so they should be read as controlled outcome proxies for mechanism-controller comparison rather than direct measurements of live-market welfare
- the response surface is still low-order and pairwise; it supports structured stress analysis, but not a fully general causal model of market behavior
- the agent behaviors are stylized and do not yet include richer strategic adaptation; their role is repeatability and controlled heterogeneity rather than faithful modeling of every strategic response

10 Related Work

The benchmark is motivated directly by frequent-batch-auction market design [3], but the paper's framing is closer to ML benchmark design and evaluation-protocol work than to a pure mechanism-design paper. The closest domain-specific relatives are replay-centric financial RL suites such as FinRL-Meta [6], multi-agent simulators such as ABIDES-Gym [1], and high-throughput LOB simulators such as JAX-LOB [4]. Relative to those systems, the present contribution is aimed at a narrower gap: controller comparisons that survive explicit settlement admissibility, mechanism variation, and welfare-sensitive measurement in the same loop.

The paper also connects to a broader ML discussion about benchmark design, constrained decision making, and objective misspecification. The main claim is not that market learning needs

another simulator per se. It is that simplifying the evaluation protocol can change method rankings and hide learning failures that only appear once admissibility and welfare-sensitive outcomes are enforced. In that sense, the benchmark is meant as a domain-specific instance of a more general evaluation problem.

11 Conclusion

This paper makes one central claim: matching-only evaluation can mislead conclusions about market-learning systems. Once settlement admissibility and welfare-sensitive outcomes are included, controller rankings and operating points change materially. The strongest evidence is the protocol-level rank shift in Figure 2: the simplified `matching_only` view does not preserve the ordering of the full benchmark. Across base runs, calibrated reruns, and counterfactual controls, controllers that optimize aggressively for latency and fill throughput tend to widen surplus-transfer gap, while value-style learners remain more stable on the latency-welfare frontier than online learners that drift toward latency-favoring operating points. The benchmark is therefore useful not only because it models more of the market stack, but because it exposes when simplified evaluation would tell the wrong story.

A Appendix Tables

Scenario	Orders/s	p99 (ms)	Arb Profit	Risk Rejects
Ablation-Control	1190.48	320.00	685.25	2922
Ablation-RelaxedRisk	1770.24	302.50	2198.50	0
Ablation-RandomTieBreak	1190.48	402.50	738.50	2908
Ablation-NoSettlementChecks	1190.48	320.00	685.25	2922

Table 11: Mechanism ablation appendix table.

Controller	p99 (ms)	Impact	Retail Surplus	Welfare Gap
BurstAware	400.00	5.31	-0.2258	0.7579
LinUCB	155.00	5.90	0.0795	2.1694
TinyMLP	221.25	5.22	-0.3128	1.9719
OfflineContextual	215.00	4.94	-0.1090	1.3769
FittedQ	145.00	5.88	0.0742	2.2036
OnlineDQN	145.00	5.92	0.0662	2.2472

Table 12: Controller comparison appendix table.

References

[1] Ibrahim Amrouni et al. Abides-gym: Gym environments for multi-agent discrete event simulation and application to financial markets, 2021.

Arb	Retail	Informed	Maker	Welfare Gap
0	1	2	1	0.2376
3	1	2	1	1.2839
0	3	2	1	0.1335
3	3	2	3	1.8005

Table 13: Selected cells from the unified arbitrage / retail / informed / maker hypercube.

[2] Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. Openai gym, 2016.

[3] Eric Budish, Peter Cramton, and John Shim. The high-frequency trading arms race: Frequent batch auctions as a market design response. *The Quarterly Journal of Economics*, 130(4):1547–1621, 2015.

[4] Felix Frey et al. Jax-lob: A gpu-accelerated limit order book simulator to unlock large scale reinforcement learning for trading, 2023.

[5] Ilya Kostrikov, Ashvin Nair, and Sergey Levine. Offline reinforcement learning with implicit q-learning, 2022.

[6] Xiao-Yang Liu et al. Finrl-meta: Market environments and benchmarks for data-driven financial reinforcement learning, 2022.

[7] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms, 2017.